



1.4.1.3.3 Letta ADE

1.4.1.3.4 1.4.1.3.4 Letta Docs

1.4.1.4 1.4.1.4 ++ Agentic

1.4.1.4.1 1.4.1.4.1 Inngest

1.4.1.4.2 Inngest Agent Kit

1.4.1.4.2.1 Agent Kit Reddit Search BrowserBase

1.4.1.4.3 No Code / Low Code

1.4.1.4.4 Flowize AI

1.4.1.4.5 Langflow

1.4.1.4.6 1.4.1.4.6 Scrape & Crawl

1.4.1.4.6.1 LightRAG

1.4.2 1.4.2 Browser Control

1.4.2.1 Browser Use

1.5 1.5 Autonomous

1.5.1 Browser User

1.5.2 Proxy Lite

1.5.3 Open Manus

1.5.4 Open Operator

1.5.5 Steel

1.5.6 BrowserBase

1.6 1.6 Production

1.6.1 Cloudflare

1.6.1.1 1.6.1.1 DNS

1.6.1.2 Security / WAF

1.7 1.7 Stack As TODOs

1.7.1 (Inside) DevAgents & DevTools

1.7.1.1 System Architecture

1.8 Database

1.8.1 SpacetimeDB

1.8.2 Neon & Turso

1.8.2.1 Neon

1.8.3 Turso

1.8.4 Memory

1.8.5 1.8.5 GenUI

1.8.5.1 Crayon

## 1.9 Infra & IaC

### 1.10 DevSecOps & CI/CD

#### 1.10.1 Docker

##### 1.10.1.1 MCP Catalog & Toolkit

##### 1.10.1.2 GH Copilot

##### 1.10.1.3 Docker Compose

##### 1.10.1.4 Docker Bake

##### 1.10.1.4.1 Testcontainers Cloud

##### 1.10.1.5 k8s

##### 1.10.1.6 Dagger

##### 1.10.1.7 Dagger Agents

##### 1.10.1.8 Uffizzi

##### 1.10.1.8.1 Temporal

##### 1.10.1.8.2 GitHub Actions

##### 1.10.1.8.3 GitHub Dev Program

### 1.11 LLMetry & Ops

#### 1.11.1 Promptfoo

#### 1.11.2 Comet / Opik

##### 1.11.2.1 Opik MCP

##### 1.11.2.2 Opik

##### 1.11.2.3 Opik GitHub

##### 1.11.2.4 Comet Docs

##### 1.11.2.5 Comet GitHub

#### 1.11.3 AgentOps

#### 1.11.4 Agenta

##### 1.11.4.1 Langwatch

### 1.12 Service Mesh

#### 1.12.1 Kong API & AI

##### 1.12.1.1 Kuma

##### 1.12.1.2 Kong Gateway

##### 1.12.1.3 Portkey

##### 1.12.1.4 Not Diamond

### 1.13 LiteLLM

### 1.14 CMS

#### 1.14.1 Prismic

## 1.15 More Agentic Frameworks

### 1.15.1 Dagger

### 1.15.2 Langbase ICP Insight Team Blogger

### 1.15.3 AI Devin

## 1.16 ML Ops

### 1.16.1 1.16.1 Prime

#### 1.16.1.1 Q Learn

## 1.17 1.17 Graph Resources

---

## ▼ Penumbra

## Market Research

---

 Page

## Competitive Analysis

---

### 1 Penumbra Competitor Analysis:

#### 1.1 AI-Powered Knowledge Management Applications

### 2 Competitive Landscape

### 3 Strategy and Positioning Recommendations for Penumbra

#### 3.1 1. Leverage Differentiation Opportunities

### 4 Recommendation:4

#### 4.1 2. Address Common Competitor Weaknesses

##### 4.1.1 Several pain points emerge across competitors:

### 4.2 Recommendation:4.2

#### 4.3 3. Hybrid Pricing Strategy

##### 4.3.1 Pricing is a key competitive factor in this space:

### 4.4 Recommendation:4.4

#### 4.5 4. Focus on Integration Ecosystem

##### 4.5.1 The ability to connect with existing tools is critical for knowledge management:

### 4.6 Recommendation:4.6

#### 4.7 5. Balanced UI/UX Approach

##### 4.7.1 User experience varies significantly across competitors:

### 4.8 Recommendation:4.8

4.9 6. Emphasize Privacy and Security

4.9.1 Enterprise concerns about AI and data security present an opportunity:

4.10 Recommendation:4.10

4.11 7. Vertical-Specific Solutions

4.11.1 Most competitors offer horizontal solutions:

4.12 Recommendation:4.12

4.13 8. Community-Driven Development

4.13.1 Build an engaged user community:

4.14 Recommendation:4.14

4.15 9. Knowledge Graph Visualization

4.15.1 A feature missing from most competitors:

4.15.2 Recommendation:4.15.2

4.16 10. AI-First, But Human-Centered

4.16.1 Emphasize the balance of AI power with human needs:

4.17 Recommendation:4.17

5 Conclusion

Findings about competitors for Penumbra as a comprehensive comparison.

Penumbra Competitor Analysis:

AI-Powered Knowledge Management Applications

Competitive Landscape

Competitor Name	Rating	Key Strengths	Key Weaknesses	Additional Information
Notion AI	4.7/5	- Highly customizable workspace- Versatile for notes, tasks, and knowledge bases- Strong AI-assisted content creation- Excellent team collaboration features- Affordable pricing with free tier	- Search functionality limitations- Performance issues with large datasets- Mobile app usability concerns- Sometimes complex learning curve	- \$8/month per user for AI features- Widely adopted across industries- Continuous feature improvements- Strong integrated ecosystem
Mem.ai	4.5/5	- Self-organizing note system-	- Limited integrations with	- Starting at \$8.33/month-

		Excellent AI-powered search- Seamless knowledge connections- Fast and intuitive interface- Strong natural language processing	other tools- Still evolving platform with changes- Higher price point for full features- Limited customization options	Strong focus on personal knowledge management- Particularly good for entrepreneurs and executives- AI-first approach to information retrieval
<b>Guru</b>	4.5/5	- Chrome extension for anywhere access- Strong verification workflows- Enterprise-grade knowledge sharing- Excellent integrations (Slack, Teams)- AI-driven search capabilities	- Steep learning curve for new users- Higher price point for full features- Requires continuous maintenance- Limited free tier functionality	- All-in-one platform- \$18/month- High satisfaction for knowledge management (G2)- 98% satisfaction rate- 50% faster implementation than competitors- Strong focus on team knowledge sharing
<b>Qatalog</b>	4.3/5	- Real-time data access without indexing- Strong security and compliance features- Search across structured/unstructured data- Permissions-based access controls- \$26M in venture funding	- Relatively newer to market- Some users report confusion about functionality- Less established user community- Limited templates and resources	- Pro plan at \$15/month per user- 14-day free trial- Strong enterprise focus- End-to-end encryption- ISO 27001, SOC2-II, GDPR compliant
<b>Saner.ai</b>	4.2/5	- Designed specifically for ADHD users- AI tag suggestion for organization- Chrome extension for research- Clean, focused interface- Natural language search	- Occasional technical issues- Requires internet access- Less established in enterprise settings- More limited feature set than competitors	- Starting at \$7.89/month- Growing user base- Strong focus on individual productivity- Easy by-side research and note taking- Position as niche tool for specific use cases

## Strategy and Positioning Recommendations for Penumbra

### 1. Leverage Differentiation Opportunities

Based on the competitive analysis, Penumbra has several opportunities to differentiate in the AI knowledge management market:

## **Recommendation:**

Position Penumbra as the most intelligent AI agent that truly understands context and user intent, going beyond simple keyword search or basic AI assistance.

- Focus on how Penumbra can anticipate user needs before they even ask.

## **2. Address Common Competitor Weaknesses**

Several pain points emerge across competitors:

### **Recommendation:**

Develop Penumbra with exceptional search capabilities that work flawlessly even with large datasets.

- Emphasize performance optimization and ensure the platform remains fast and responsive regardless of knowledge base size.

## **3. Hybrid Pricing Strategy**

Pricing is a key competitive factor in this space:

### **Recommendation:**

Offer a compelling free tier with meaningful functionality (unlike Guru's limited free version) to drive adoption, with a competitively priced premium tier around \$10/month (positioning between Saner.ai and Notion, but below Guru and Qatalog).

- Consider a special "Founding Member" pricing for early adopters.

## **4. Focus on Integration Ecosystem**

The ability to connect with existing tools is critical for knowledge management:

### **Recommendation:**

Develop a robust API and integration ecosystem from day one, ensuring Penumbra works seamlessly with tools like Slack, Microsoft Teams, Google Workspace, and other productivity apps.

- Create a visual "knowledge graph" that shows connections between information across various platforms.

## **5. Balanced UI/UX Approach**

User experience varies significantly across competitors:

## **Recommendation:**

Strike the perfect balance between Notion's flexibility and Mem.ai's simplicity.

- Create an interface that adapts to the user's experience level, offering both guided experiences for beginners and advanced capabilities for power users.

## **6. Emphasize Privacy and Security**

Enterprise concerns about AI and data security present an opportunity:

### **Recommendation:**

Position Penumbra as the most secure and privacy-focused knowledge management solution.

Highlight how Penumbra can work with sensitive information without requiring data to leave the company's environment (similar to Qatalog's approach, but with more emphasis).

## **7. Vertical-Specific Solutions**

Most competitors offer horizontal solutions:

### **Recommendation:**

Develop specialized versions of Penumbra for specific industries like legal, healthcare, education, and software development with pre-built templates, workflows, and AI agents trained on industry-specific knowledge.

## **8. Community-Driven Development**

Build an engaged user community:

### **Recommendation:**

Create a vibrant community around Penumbra with user forums, template sharing, regular webinars, and a transparent product roadmap where users can vote on features.

- Position Penumbra as the tool that truly listens to its users.

## **9. Knowledge Graph Visualization**

A feature missing from most competitors:

### **Recommendation:**

Develop an innovative visualization interface that allows users to see connections between their knowledge in ways that other tools cannot, helping users discover insights they wouldn't otherwise find.

## **10. AI-First, But Human-Centered**



Emphasize the balance of AI power with human needs:

## Recommendation:

Position Penumbra as "AI-powered, human-centered" – the perfect balance of cutting-edge AI with thoughtful, human-friendly design.

- Focus on how Penumbra enhances human creativity and productivity rather than replacing human thinking.


## Conclusion

The AI-powered knowledge management space is competitive but still evolving rapidly. By positioning Penumbra as a next-generation solution that addresses the core weaknesses of current offerings while introducing innovative features, there is a significant opportunity to capture market share. The key will be delivering a product that balances powerful AI capabilities with intuitive usability, strong privacy features, and compelling pricing.

---

## Product / UX

### ▼ Example PRD → XPRT

 PRD: XPRT

 [Page](#)

## Requirements Gathering

---

 [Page](#)

## Requirements From Notion

---

1 7 Primitives

2 v5 beta prompt

2.1 !Important

2.2 Your Role and Capabilities

2.3 Key Concepts in Penumbra

2.3.1 Users

2.3.2 Nodes

2.3.3 Edges

2.3.4 Schemas

2.4 Available Tools

- 2.4.1 Worldview Components
- 2.4.2 Founder Worldview Management
- 2.4.3 Belief Management
- 2.4.4 Narrative Management
- 2.4.5 Operating Constraint Management
- 2.4.6 Strategic Elements
- 2.4.7 Strategic Hypothesis Management
- 2.4.8 Strategic Initiative Management
- 2.4.9 Internal Objective Management
- 2.4.10 Actor Management
- 2.4.11 Actor Management
- 2.4.12 Thinking Infrastructure
- 2.4.13 Concept Management
- 2.4.14 Supporting Elements
- 2.4.15 Document Management
- 2.4.16 Industry Classification
- 2.4.17 Generic Operations
- 2.4.18 Node Operations
- 2.4.19 Edge Operations
- 2.4.20 Metadata Operations
- 2.4.21 Query Operations

## 2.5 Best Practices

- 2.5.1 Belief Definition
- 2.5.2 Narrative Construction
- 2.5.3 Strategic Reasoning
- 2.5.4 Actor Management
- 2.5.5 Relationship Management
- 2.5.6 Metadata Usage

## 3 v5 beta prompt

- 3.1 Important
- 3.2 Your Role and Capabilities
- 3.3 Key Concepts in Penumbra
  - 3.3.1 Users
  - 3.3.2 Nodes
  - 3.3.3 Edges

#### 3.3.4 Schemas

### 3.4 Available Tools

#### 3.4.1 Worldview Components

#### 3.4.2 Founder Worldview Management

#### 3.4.3 Belief Management

#### 3.4.4 Narrative Management

#### 3.4.5 Operating Constraint Management

#### 3.4.6 Strategic Elements

#### 3.4.7 Strategic Hypothesis Management

#### 3.4.8 Strategic Initiative Management

#### 3.4.9 Internal Objective Management

#### 3.4.10 Actor Management

#### 3.4.11 Actor Management

#### 3.4.12 Thinking Infrastructure

#### 3.4.13 Concept Management

#### 3.4.14 Supporting Elements

#### 3.4.15 Document Management

#### 3.4.16 Industry Classification

#### 3.4.17 Generic Operations

#### 3.4.18 Node Operations

#### 3.4.19 Edge Operations

#### 3.4.20 Metadata Operations

#### 3.4.21 Query Operations

### 3.5 Best Practices

#### 3.5.1 Belief Definition

#### 3.5.2 Narrative Construction

#### 3.5.3 Strategic Reasoning

#### 3.5.4 Actor Management

#### 3.5.5 Relationship Management

#### 3.5.6 Metadata Usage

---

## 7 Primitives

### v5 beta prompt

You are a helpful assistant who is helping the user build a worldview graph – a specialized knowledge graph that models how the user thinks and strategizes. You help users externalize, structure, and leverage their worldview into a semantic knowledge graph using the tools available to you.

## !Important

When the user initiates a chat, start by using the `read_graph` tool to pull the contents of their graph into your context window.

## Your Role and Capabilities

As a Penumbra Assistant, you can use your available tools and critical thinking skills to:

- Call the available tools to structure the founder's worldview and its supporting elements into a coherent and connected graph
- Reference the user's graph to in order to help them make decisions aligned with their worldview
- Help users discover insights and potential contradictions in their thinking

## Key Concepts in Penumbra

### Users

Users in Penumbra are primarily startup founders, business leaders, and strategic consultants.

### Nodes

Nodes represent entities in the founder's cognitive and strategic system:

- **Actor:** Humans, AI agents, or organizational units with agency
- **Belief:** Propositional claims held to be true
- **Concept:** Terms or ideas in the founder's specific ontology
- **Document:** Sources, evidence, or artifacts
- **Founder Worldview:** The cognitive root of the founder or organization
- **Industry:** Business sectors providing market context
- **Internal Objective:** Specific goals or requirements within initiatives
- **Narrative:** Structured expressions of the worldview for different audiences
- **Operating Constraint:** Hard boundaries that define limits
- **Strategic Hypothesis:** Testable assumptions derived from beliefs
- **Strategic Initiative:** Concrete efforts that test hypotheses

### Edges

Edges represent relationships between entities, organized by category:

- **Structural:** HAS\_BELIEF, PART\_OF\_WORLDVIEW, MEMBER\_OF, INITIATIVE\_OF\_WORLDVIEW
- **Causal:** LEADS\_TO, CAUSES, PREVENTS
- **Dependency:** DEPENDS\_ON, REQUIRES, ENABLES
- **Evidential:** JUSTIFIED\_BY, SUPPORTS, CONTRADICTS, EVIDENCES\_BELIEF
- **Associative:** RELATED\_TO, SIMILAR\_TO, RELATED\_TO\_BELIEF
- **Temporal:** PRECEDES, FOLLOWS
- **Agential:** CREATED\_BY, MAINTAINED\_BY, UNDERSTOOD\_BY, OWNED\_BY
- **Cognitive:** BELIEVES, CHALLENGES, EXPRESSES\_BELIEF, UNDERPINS\_HYPOTHESIS

## Schemas

Schemas define the structure and properties of each node type, ensuring coherent representation of the founder's worldview.

---

## Available Tools

### Worldview Components

#### Founder Worldview Management

- `add_founder_worldview`: Define the central worldview
- `update_founder_worldview`: Update the worldview
- `delete_founder_worldview`: Remove a worldview

#### Belief Management

- `add_belief`: Add a foundational belief
- `update_belief`: Update a belief
- `delete_belief`: Remove a belief

#### Narrative Management

- `add_narrative`: Define a narrative
- `update_narrative`: Update a narrative
- `delete_narrative`: Remove a narrative

#### Operating Constraint Management

- `add_operating_constraint`: Define a boundary
- `update_operating_constraint`: Update a constraint
- `delete_operating_constraint`: Remove a constraint

## Strategic Elements

### Strategic Hypothesis Management

- `add_strategic_hypothesis`: Create a hypothesis
- `update_strategic_hypothesis`: Update a hypothesis
- `delete_strategic_hypothesis`: Remove a hypothesis

### Strategic Initiative Management

- `add_strategic_initiative`: Define an initiative
- `update_strategic_initiative`: Update an initiative
- `delete_strategic_initiative`: Remove an initiative

### Internal Objective Management

- `add_internal_objective`: Define an objective
- `update_internal_objective`: Update an objective
- `delete_internal_objective`: Remove an objective

## Actor Management

### Actor Management

- `add_actor`: Add a human, AI agent, or org unit
- `update_actor`: Update an actor
- `delete_actor`: Remove an actor

## Thinking Infrastructure

### Concept Management

- `add_concept`: Define a concept
- `update_concept`: Update a concept
- `delete_concept`: Remove a concept

## Supporting Elements

### Document Management

- `add_document`: Add a document
- `update_document`: Update a document
- `delete_document`: Remove a document

## Industry Classification

- `add_industry` : Add an industry context
- `update_industry` : Update an industry
- `delete_industry` : Remove an industry

## Generic Operations

If the user has a large batch of actions to be performed, you can batch node and edge operations as long as you follow the properly defined schemas for each respective operation.

## Node Operations

- `add_nodes(nodes)` : Add multiple nodes at once
- `update_nodes(nodes)` : Update multiple nodes at once
- `delete_nodes(nodeNames)` : Delete multiple nodes at once

## Edge Operations

- `add_edges(edges)` : Add relationships between nodes
- `update_edges(edges)` : Update existing relationships
- `delete_edges(edges)` : Delete relationships

## Metadata Operations

- `add_metadata(metadata)` : Add metadata to nodes
- `delete_metadata(deletions)` : Remove metadata from nodes

## Query Operations

- `read_graph()` : Retrieve the entire knowledge graph
  - `search_nodes(query)` : Search for nodes matching a query
  - `open_nodes(names)` : Retrieve specific nodes by name
- 

## Best Practices

- Connect all items using appropriate edges / relationships to form a coherent network

## Belief Definition

- Make beliefs as specific and falsifiable as possible
- Distinguish between different types of beliefs (epistemic, strategic, etc.)
- Assign appropriate certainty levels based on evidence
- Link beliefs to their justifications and evidence

## Narrative Construction

- Target narratives to specific audiences (team, investors, customers)
- Ensure narratives express but don't contradict beliefs
- Check alignment between different narratives
- Update narratives as beliefs evolve
- Use narratives to make abstract beliefs concrete

## **Strategic Reasoning**

- Form hypotheses that are testable versions of beliefs
- Design initiatives specifically to test hypotheses
- Set clear success criteria for hypotheses
- Ensure initiatives connect back to the core worldview
- Close the loop by updating beliefs based on outcomes

## **Actor Management**

- Distinguish between human actors, AI agents, and organizational units
- Track belief alignment between actors
- Assign clear ownership for initiatives and hypotheses
- Document when actors initiate reflections or checkpoints
- Model how actors' beliefs influence worldview evolution

## **Relationship Management**

- Use specific relationship types from the registry
- Ensure relationships are useful to the worldview model
- Include relationship categories for clarity (structural, causal, etc.)

## **Metadata Usage**

- Use metadata for attributes that don't warrant their own nodes
  - Keep metadata concise and relevant
  - Prefer explicit relationships over generic metadata
- 

## **v5 beta prompt**

You are a helpful assistant who is helping the user build a worldview graph – a specialized knowledge graph that models how the user thinks and strategizes. You help users externalize, structure, and leverage their worldview into a semantic knowledge graph using the tools available to you.

## **!Important**



When the user initiates a chat, start by using the read\_graph tool to pull the contents of their graph into your context window.

## Your Role and Capabilities

As a Penumbra Assistant, you can use your available tools and critical thinking skills to:

- Call the available tools to structure the founder's worldview and its supporting elements into a coherent and connected graph
- Reference the user's graph to in order to help them make decisions aligned with their worldview
- Help users discover insights and potential contradictions in their thinking

## Key Concepts in Penumbra

### Users

Users in Penumbra are primarily startup founders, business leaders, and strategic consultants.

### Nodes

Nodes represent entities in the founder's cognitive and strategic system:

- **Actor:** Humans, AI agents, or organizational units with agency
- **Belief:** Propositional claims held to be true
- **Concept:** Terms or ideas in the founder's specific ontology
- **Document:** Sources, evidence, or artifacts
- **Founder Worldview:** The cognitive root of the founder or organization
- **Industry:** Business sectors providing market context
- **Internal Objective:** Specific goals or requirements within initiatives
- **Narrative:** Structured expressions of the worldview for different audiences
- **Operating Constraint:** Hard boundaries that define limits
- **Strategic Hypothesis:** Testable assumptions derived from beliefs
- **Strategic Initiative:** Concrete efforts that test hypotheses

### Edges

Edges represent relationships between entities, organized by category:

- **Structural:** HAS\_BELIEF, PART\_OF\_WORLDVIEW, MEMBER\_OF, INITIATIVE\_OF\_WORLDVIEW
- **Causal:** LEADS\_TO, CAUSES, PREVENTS
- **Dependency:** DEPENDS\_ON, REQUIRES, ENABLES
- **Evidential:** JUSTIFIED\_BY, SUPPORTS, CONTRADICTS, EVIDENCES\_BELIEF
- **Associative:** RELATED\_TO, SIMILAR\_TO, RELATED\_TO\_BELIEF

- **Temporal:** PRECEDES, FOLLOWS
- **Agential:** CREATED\_BY, MAINTAINED\_BY, UNDERSTOOD\_BY, OWNED\_BY
- **Cognitive:** BELIEVES, CHALLENGES, EXPRESSES\_BELIEF, UNDERPINS\_HYPOTHESIS

## Schemas

Schemas define the structure and properties of each node type, ensuring coherent representation of the founder's worldview.

## Available Tools

### Worldview Components

#### Founder Worldview Management

- `add_founder_worldview`: Define the central worldview
- `update_founder_worldview`: Update the worldview
- `delete_founder_worldview`: Remove a worldview

#### Belief Management

- `add_belief`: Add a foundational belief
- `update_belief`: Update a belief
- `delete_belief`: Remove a belief

#### Narrative Management

- `add_narrative`: Define a narrative
- `update_narrative`: Update a narrative
- `delete_narrative`: Remove a narrative

#### Operating Constraint Management

- `add_operating_constraint`: Define a boundary
- `update_operating_constraint`: Update a constraint
- `delete_operating_constraint`: Remove a constraint

### Strategic Elements

#### Strategic Hypothesis Management

- `add_strategic_hypothesis`: Create a hypothesis
- `update_strategic_hypothesis`: Update a hypothesis
- `delete_strategic_hypothesis`: Remove a hypothesis

## Strategic Initiative Management

- `add_strategic_initiative`: Define an initiative
- `update_strategic_initiative`: Update an initiative
- `delete_strategic_initiative`: Remove an initiative

## Internal Objective Management

- `add_internal_objective`: Define an objective
- `update_internal_objective`: Update an objective
- `delete_internal_objective`: Remove an objective

## Actor Management

### Actor Management

- `add_actor`: Add a human, AI agent, or org unit
- `update_actor`: Update an actor
- `delete_actor`: Remove an actor

## Thinking Infrastructure

### Concept Management

- `add_concept`: Define a concept
- `update_concept`: Update a concept
- `delete_concept`: Remove a concept

## Supporting Elements

### Document Management

- `add_document`: Add a document
- `update_document`: Update a document
- `delete_document`: Remove a document

### Industry Classification

- `add_industry`: Add an industry context
- `update_industry`: Update an industry
- `delete_industry`: Remove an industry

## Generic Operations

If the user has a large batch of actions to be performed, you can batch node and edge operations as long as you follow the properly defined schemas for each respective operation.

## Node Operations

- `add_nodes(nodes)` : Add multiple nodes at once
- `update_nodes(nodes)` : Update multiple nodes at once
- `delete_nodes(nodeNames)` : Delete multiple nodes at once

## Edge Operations

- `add_edges(edges)` : Add relationships between nodes
- `update_edges(edges)` : Update existing relationships
- `delete_edges(edges)` : Delete relationships

## Metadata Operations

- `add_metadata(metadata)` : Add metadata to nodes
- `delete_metadata(deletions)` : Remove metadata from nodes

## Query Operations

- `read_graph()` : Retrieve the entire knowledge graph
  - `search_nodes(query)` : Search for nodes matching a query
  - `open_nodes(names)` : Retrieve specific nodes by name
- 

## Best Practices

- Connect all items using appropriate edges / relationships to form a coherent network

## Belief Definition

- Make beliefs as specific and falsifiable as possible
- Distinguish between different types of beliefs (epistemic, strategic, etc.)
- Assign appropriate certainty levels based on evidence
- Link beliefs to their justifications and evidence

## Narrative Construction

- Target narratives to specific audiences (team, investors, customers)
- Ensure narratives express but don't contradict beliefs
- Check alignment between different narratives
- Update narratives as beliefs evolve
- Use narratives to make abstract beliefs concrete

## Strategic Reasoning

- Form hypotheses that are testable versions of beliefs
- Design initiatives specifically to test hypotheses
- Set clear success criteria for hypotheses
- Ensure initiatives connect back to the core worldview
- Close the loop by updating beliefs based on outcomes

## Actor Management

- Distinguish between human actors, AI agents, and organizational units
- Track belief alignment between actors
- Assign clear ownership for initiatives and hypotheses
- Document when actors initiate reflections or checkpoints
- Model how actors' beliefs influence worldview evolution

## Relationship Management

- Use specific relationship types from the registry
- Ensure relationships are useful to the worldview model
- Include relationship categories for clarity (structural, causal, etc.)

## Metadata Usage

- Use metadata for attributes that don't warrant their own nodes
- Keep metadata concise and relevant
- Prefer explicit relationships over generic metadata

---

### ▼ Planning

#### ▼ Database

---

 Page

## Data Modeling

---

### 1 Optimal Data Modeling Approach for Penumbra

#### 1.1 1. Domain-Driven Separation of Concerns

##### 1.1.1 PostgreSQL (Relational Data)

###### 1.1.1.1 • 1.1.1.1 User and Authentication Data

###### 1.1.1.2 • 1.1.1.2 Operational/Transactional Data

### 1.1.2 Neo4j (Graph Data)

#### 1.1.2.1 • 1.1.2.1 Knowledge Graph Entities

#### 1.1.2.2 • 1.1.2.2 Semantic Relationships

### 1.2 2. Cross-Database Reference Architecture

#### 1.2.1 Identifier Strategy

#### 1.2.2 Example pattern:

### 1.3 Synchronization Patterns

#### 1.3.1 • 1.3.1 Event-Driven Consistency

#### 1.3.1.1 • 1.3.1.1 Example flow:

### 1.4 3. Schema Design Patterns

#### 1.4.1 PostgreSQL Schema

#### 1.4.2 Neo4j Graph Model (or MemGraph)

### 1.5 4. Data Access Layer

#### 1.5.1 Repository Pattern Implementation

### 1.6 5. Query Optimization Strategies

#### 1.6.1 PostgreSQL Optimization

#### 1.6.1.1 • 1.6.1.1 Example

#### 1.6.2 Neo4j Optimization

#### 1.6.2.1 • 1.6.2.1 Example

### 1.7 6. Integration with Vector Embeddings

#### 1.7.1 Hybrid Search Architecture

### 1.8 7. Cognitive Shapes Data Model

#### 1.8.1 Core Shape Definitions

#### 1.8.2 PostgreSQL Storage for Shapes

#### 1.8.3 Neo4j for Runtime Representation (or MemGraph)

### 1.9 8. Data Flow and Transaction Management

#### 1.9.1 Atomic Operations

### 1.10 9. API Layer Design

#### 1.10.1 GraphQL API for Unified Access

### 1.11 10. Caching Strategy

### 1.12 11. Recommendations for Penumbra Implementation

#### 1.12.1 1. 1.12.1 Start with Clear Domain Boundaries

#### 1.12.2 2. 1.12.2 Implement a Service Layer Abstraction

#### 1.12.3 3. 1.12.3 Build Resilient Cross-Database Operations

#### 1.12.4 4. 1.12.4 Optimize for Query Patterns

1.12.5 5. 1.12.5 Use Composite Identifiers Consistently

1.12.6 6. 1.12.6 Implement Monitoring for Both Databases

1.12.7 7. 1.12.7 Design for Scale from the Start

1.12.8 8. 1.12.8 Leverage Database-Specific Features

1.13 Local KG Example (From my Dev Environment)

## 2 Knowledge Graph Visualization

2.1 Diagram Explanation

2.2 9. 2.2 Node Types2.2 :

2.2.1 • 2.2.1 System nodes (pink):

2.2.1.1 • 2.2.1.1 LLM and Knowledge Graph

2.2.2 • 2.2.2 Capability nodes (blue):

2.2.2.1 • 2.2.2.1 Natural Language Understanding and Schema Inference

2.3 10. 2.3 Relationships2.3 :

2.3.1 • 2.3.1 Direct relationships shown with solid lines

2.3.2 • 2.3.2 Observational properties shown with dotted lines to text boxes

2.4 11. 2.4 Properties2.4 :

2.4.1 • 2.4.1 Each node's key properties are listed in associated text boxes

2.4.2 • 2.4.2 Properties help explain the role and capabilities of each component

2.5 12. 2.5 Visual Hierarchy2.5 :

2.5.1 • 2.5.1 Central nodes represent core systems

2.5.2 • 2.5.2 Supporting capabilities branch out

2.5.3 • 2.5.3 Observations provide additional context

---

# Optimal Data Modeling Approach for Penumbra

## 1. Domain-Driven Separation of Concerns

### PostgreSQL (Relational Data)

- **User and Authentication Data**
  - User profiles, credentials, access permissions
  - Organization membership and hierarchies
  - API keys and security tokens
- **Operational/Transactional Data**

- Chat sessions and message metadata
- Document metadata and file references
- Usage statistics and analytics
- Audit logs and system events

## Neo4j (Graph Data)

- **Knowledge Graph Entities**
  - Core cognitive entities (Beliefs, Principles, Strategies)
  - Domain-specific concepts and their relationships
  - Hierarchical and networked knowledge structures
- **Semantic Relationships**
  - Entity connections with typed relationships
  - Weighted/scored relationships
  - Temporal and contextual relationship metadata

## 2. Cross-Database Reference Architecture

### Identifier Strategy

- Use a consistent UUID strategy across both databases
- Maintain a reference mapping layer for entities that exist in both systems

### Example pattern:

```

TypeScript
// PostgreSQL entity with graph reference
interface Document {
  id: string;           // UUID primary key
  title: string;
  content: string;
  graphNodeIds: string[]; // References to Neo4j nodes
}

// Neo4j node with relational reference
interface KnowledgeNode {
  id: string;           // Same UUID format
  type: string;         // Node type (Belief, Strategy, etc.)
  content: string;      // Core content
  metadata: {
    postgresId?: string; // Optional reference back to PostgreSQL
  }
}

```



## Synchronization Patterns

- **Event-Driven Consistency**

- Use an event bus (e.g., lightweight message queue) for cross-database updates
- Implement the Outbox Pattern for reliable cross-database transactions
- **Example flow:**
  1. Write to primary database
  2. Store outbox event for secondary database update
  3. Process outbox asynchronously
  4. Confirm consistency through reconciliation processes

## 3. Schema Design Patterns

### PostgreSQL Schema

SQL

```

-- Core user data
CREATE TABLE users (
  id UUID PRIMARY KEY,
  email TEXT UNIQUE NOT NULL,
  name TEXT,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
  updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

-- Organizations
CREATE TABLE organizations (
  id UUID PRIMARY KEY,
  name TEXT NOT NULL,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
  updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

-- Organization membership
CREATE TABLE organization_members (
  id UUID PRIMARY KEY,
  user_id UUID REFERENCES users(id),
  organization_id UUID REFERENCES organizations(id),
  role TEXT NOT NULL,
  UNIQUE(user_id, organization_id)
);

-- Chat sessions
CREATE TABLE chat_sessions (
  id UUID PRIMARY KEY,
  user_id UUID REFERENCES users(id),
  title TEXT,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
  updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

-- Chat messages with graph references
CREATE TABLE chat_messages (
  id UUID PRIMARY KEY,
  session_id UUID REFERENCES chat_sessions(id),
  content TEXT NOT NULL,
  role TEXT NOT NULL,
  graph_node_references JSONB, -- Store references to Neo4j nodes
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

-- Documents/artifacts with graph metadata
CREATE TABLE documents (
  id UUID PRIMARY KEY,
  user_id UUID REFERENCES users(id),
  title TEXT NOT NULL,
  content TEXT,

```

```

file_path TEXT,
mime_type TEXT,
graph_node_references JSONB, -- Store references to Neo4j nodes
created_at TIMESTAMPT WITH TIME ZONE DEFAULT NOW(),
updated_at TIMESTAMPT WITH TIME ZONE DEFAULT NOW()
);

```

## Neo4j Graph Model (or MemGraph)

```

SQL

// Core node types
CREATE (n:Belief {
  id: $id,
  content: $content,
  confidence: $confidence,
  source: $source,
  created_by: $userId,
  created_at: datetime()
})
CREATE (n:Principle {
  id: $id,
  content: $content,
  importance: $importance,
  domain: $domain,
  created_by: $userId,
  created_at: datetime()
})
CREATE (n:Strategy {
  id: $id,
  content: $content,
  timeframe: $timeframe,
  status: $status,
  created_by: $userId,
  created_at: datetime()
})

// Example relationships
CREATE (b:Belief {id: $beliefId})-[:SUPPORTS {strength: $strength}]->
(p:Principle {id: $principleId})
CREATE (p:Principle {id: $principleId})-[:INFORMS {context:
$content}]->(s:Strategy {id: $strategyId})
CREATE (s1:Strategy {id: $strategy1Id})-[:DEPENDS_ON {criticality:
$criticality}]->(s2:Strategy {id: $strategy2Id})

```

## 4. Data Access Layer

### Repository Pattern Implementation

TypeScript

```

// Base repository with common functionality
abstract class BaseRepository<T> {
    abstract findById(id: string): Promise<T | null>;
    abstract create(entity: Omit<T, 'id'>): Promise<T>;
    abstract update(id: string, entity: Partial<T>): Promise<T>;
    abstract delete(id: string): Promise<boolean>;
}

// PostgreSQL repository implementation
class PostgresRepository<T> extends BaseRepository<T> {
    constructor(private table: string, private db: DrizzleClient) {
        super();
    }

    async findById(id: string): Promise<T | null> {
        return this.db.query.from(this.table).where(eq(this.table.id, id)).first();
    }

    // Other CRUD implementations
}

// Neo4j repository implementation
class Neo4jRepository<T> extends BaseRepository<T> {
    constructor(private label: string, private session: Neo4jSession) {
        super();
    }

    async findById(id: string): Promise<T | null> {
        const result = await this.session.run(
            `MATCH (n:${this.label} {id: $id}) RETURN n`,
            { id }
        );
        return result.records.length > 0 ?
this.mapRecord(result.records[0]) : null;
    }

    // Other CRUD implementations
}

// Higher-level service that coordinates between repositories
class KnowledgeService {
    constructor(
        private userRepo: PostgresRepository<User>,
        private documentRepo: PostgresRepository<Document>,
        private beliefRepo: Neo4jRepository<Belief>,
        private principleRepo: Neo4jRepository<Principle>,
        private strategyRepo: Neo4jRepository<Strategy>,
        private graphQueryService: GraphQueryService
    ) {}
}

```

```
// Complex operations that span both databases
async getUserKnowledgeGraph(userId: string):
Promise<KnowledgeGraph> {
  const user = await this.userRepo.findById(userId);
  if (!user) throw new Error('User not found');

  const documents = await this.documentRepo.findById(userId);
  const graphData = await
this.graphQueryService.getUserGraph(userId);

  return {
    user,
    documents,
    graph: graphData
  };
}
}
```

## 5. Query Optimization Strategies

### PostgreSQL Optimization

- Use appropriate indexes on frequently queried columns
- Implement materialized views for complex aggregations
- Use JSON/JSONB for flexible schema needs within structured data
- Example

```
SQL
-- Index for user lookup by email
CREATE INDEX idx_users_email ON users(email);

-- Index for organization membership lookup
CREATE INDEX idx_org_members_user_id ON
organization_members(user_id);
CREATE INDEX idx_org_members_org_id ON
organization_members(organization_id);

-- JSONB index for graph reference lookup
CREATE INDEX idx_chat_messages_graph_refs ON chat_messages USING
gin(graph_node_references);
```

### Neo4j Optimization

- Create appropriate indexes on node properties used in queries
- Use relationship types effectively to optimize traversals
- Implement graph projections for frequently accessed subgraphs
- Example

SQL

```
// Create indexes for node lookups
CREATE INDEX idx_belief_id FOR (n:Belief) ON (n.id);
CREATE INDEX idx_principle_id FOR (n:Principle) ON (n.id);
CREATE INDEX idx_strategy_id FOR (n:Strategy) ON (n.id);

// Create composite index for user-created nodes
CREATE INDEX idx_node_user_created FOR (n) ON (n.created_by,
n.created_at);
```

## 6. Integration with Vector Embeddings

### Hybrid Search Architecture

- Store vector embeddings in both databases:
  - PostgreSQL: Use pgvector for document-level embeddings
  - Neo4j: Store embeddings as node properties for semantic similarity
- Example implementation:



TypeScript

```
// Generate embedding for a document or concept
async function generateEmbedding(text: string): Promise<number[]> {
  const embedding = await embeddingService.embed(text);
  return embedding;
}

// Store in PostgreSQL with pgvector
async function storeDocumentWithEmbedding(doc: Document) {
  const embedding = await generateEmbedding(doc.content);
  await db.execute(
    `INSERT INTO documents(id, title, content, embedding)
    VALUES($1, $2, $3, $4)`,
    [doc.id, doc.title, doc.content, embedding]
  );
}

// Store in Neo4j
async function storeConceptWithEmbedding(concept: Concept) {
  const embedding = await generateEmbedding(concept.content);
  await session.run(
    `CREATE (c:Concept {
      id: $id,
      content: $content,
      embedding: $embedding
    })`,
    {
      id: concept.id,
      content: concept.content,
      embedding: embedding
    }
  );
}
```

## 7. Cognitive Shapes Data Model

Implement a modular pattern for the Cognitive Shapes architecture:

### Core Shape Definitions

```

TypeScript

// Core shape definition
interface CognitiveShape {
  id: string;
  name: string;
  type: string;
  description: string;
  capabilities: string[];
  tools: Tool[];
  resources: Resource[];
  prompts: Prompt[];
}

// Tool definition
interface Tool {
  id: string;
  name: string;
  description: string;
  type: 'action' | 'query' | 'reasoning';
  handler: string; // Reference to function implementation
  parameters: Parameter[];
  returns: ReturnType;
}

// Resource definition
interface Resource {
  id: string;
  name: string;
  type: string;
  source: string;
  accessPattern: string;
  schema?: JSONSchema;
}

```

## PostgreSQL Storage for Shapes

SQL

```
// PostgreSQL storage for shapes
CREATE TABLE cognitive_shapes (
  id UUID PRIMARY KEY,
  name TEXT NOT NULL,
  type TEXT NOT NULL,
  description TEXT,
  capabilities JSONB,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
  updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

CREATE TABLE shape_tools (
  id UUID PRIMARY KEY,
  shape_id UUID REFERENCES cognitive_shapes(id),
  name TEXT NOT NULL,
  description TEXT,
  type TEXT NOT NULL,
  handler TEXT NOT NULL,
  parameters JSONB,
  returns JSONB,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
  updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);
```

## Neo4j for Runtime Representation (or MemGraph)

SQL

```
// Neo4j representation for runtime graph traversal
CREATE (s:CognitiveShape {
  id: $id,
  name: $name,
  type: $type
})

CREATE (t:Tool {
  id: $id,
  name: $name,
  type: $type,
  handler: $handler
})

CREATE (s)-[:HAS_TOOL]->(t)
```

## 8. Data Flow and Transaction Management

### Atomic Operations

- Use database-specific transactions for operations within a single database

- Implement the Saga pattern for operations that span both databases
- Example:

TypeScript

```

async function createUserWithKnowledgeGraph(userData, initialBeliefs)
{
  // Start a coordinated transaction
  const transactionId = generateUUID();

  try {
    // PostgreSQL transaction
    const user = await db.transaction(async (tx) => {
      const newUser = await userRepo.create(userData);
      await outboxRepo.create({
        id: generateUUID(),
        transactionId,
        type: 'USER_CREATED',
        payload: newUser,
        status: 'PENDING'
      });
      return newUser;
    });

    // Neo4j transaction
    const beliefs = await neo4jSession.writeTransaction(async (tx) =>
    {
      const createdBeliefs = [];
      for (const belief of initialBeliefs) {
        const result = await tx.run(
          `CREATE (b:Belief {
            id: $id,
            content: $content,
            created_by: $userId
          }) RETURN b`,
          {
            id: generateUUID(),
            content: belief.content,
            userId: user.id
          }
        );
        createdBeliefs.push(result.records[0].get('b').properties);
      }
      return createdBeliefs;
    });

    // Mark transaction as complete
    await outboxRepo.update(
      { transactionId },
      { status: 'COMPLETED' }
    );

    return { user, beliefs };
  } catch (error) {
    // Implement compensation logic
    await outboxRepo.update(

```

```
        { transactionId },  
        { status: 'FAILED', error: error.message }  
    );  
    throw error;  
  }  
}
```

## 9. API Layer Design

### GraphQL API for Unified Access

GraphQL



```
// GraphQL schema combining both data sources
const typeDefs = gql`
  type User {
    id: ID!
    email: String!
    name: String
    organizations: [Organization!]
    beliefs: [Belief!]
    principles: [Principle!]
    strategies: [Strategy!]
  }

  type Organization {
    id: ID!
    name: String!
    members: [User!]
  }

  type Belief {
    id: ID!
    content: String!
    confidence: Float
    supportsPrinciples: [Principle!]
    createdBy: User
  }

  type Principle {
    id: ID!
    content: String!
    importance: Float
    supportedByBeliefs: [Belief!]
    informsStrategies: [Strategy!]
    createdBy: User
  }

  type Strategy {
    id: ID!
    content: String!
    timeframe: String
    status: String
    informedByPrinciples: [Principle!]
    dependsOn: [Strategy!]
    createdBy: User
  }

  type Query {
    user(id: ID!): User
    userBeliefs(userId: ID!): [Belief!]
    beliefNetwork(beliefId: ID!, depth: Int): [Belief!]
    strategyDependencies(strategyId: ID!): [Strategy!]
  }
`
```

```

type Mutation {
  createBelief(content: String!, confidence: Float): Belief
  connectBeliefToPrinciple(beliefId: ID!, principleId: ID!,
strength: Float): Boolean
  createStrategy(content: String!, timeframe: String, principleIds:
[ID!]): Strategy
}
`
;

// Resolver implementation
const resolvers = {
  Query: {
    user: async (_, { id }) => {
      return await userRepo.findById(id);
    },
    userBeliefs: async (_, { userId }) => {
      return await neo4jSession.run(
        `MATCH (b:Belief {created_by: $userId}) RETURN b`,
        { userId }
      ).then(result => result.records.map(r =>
r.get('b').properties));
    },
    // Other query resolvers
  },
  User: {
    organizations: async (user) => {
      return await orgRepo.findByUserId(user.id);
    },
    beliefs: async (user) => {
      return await neo4jSession.run(
        `MATCH (b:Belief {created_by: $userId}) RETURN b`,
        { userId: user.id }
      ).then(result => result.records.map(r =>
r.get('b').properties));
    },
    // Other user field resolvers
  },
  // Other type resolvers
};

```

## 10. Caching Strategy

Implement a multi-level caching strategy:

TypeScript

```
// In-memory cache for frequently accessed data
const nodeCache = new NodeCache({ stdTTL: 300, checkperiod: 60 });

// Redis cache for distributed caching
const redisClient = createClient({ url: process.env.REDIS_URL });

// Caching middleware for graph queries
async function cachedGraphQuery(query, params, ttl = 300) {
  const cacheKey = `graph:${query}:${JSON.stringify(params)}`;

  // Check in-memory cache first
  const cachedResult = nodeCache.get(cacheKey);
  if (cachedResult) return cachedResult;

  // Check Redis cache
  const redisCached = await redisClient.get(cacheKey);
  if (redisCached) {
    const result = JSON.parse(redisCached);
    nodeCache.set(cacheKey, result);
    return result;
  }

  // Execute query
  const result = await neo4jSession.run(query, params);
  const data = result.records.map(r => r.toObject());

  // Store in caches
  nodeCache.set(cacheKey, data);
  await redisClient.set(cacheKey, JSON.stringify(data), 'EX', ttl);

  return data;
}
```

## 11. Recommendations for Penumbra Implementation

### 1. Start with Clear Domain Boundaries

- Define precisely what belongs in PostgreSQL vs. Neo4j
- Document the cross-reference strategy between databases

### 2. Implement a Service Layer Abstraction

- Create domain-specific services that handle cross-database operations
- Abstract the underlying database complexity from application logic

### 3. Build Resilient Cross-Database Operations

- Use the Outbox pattern for reliable cross-database updates

- Implement reconciliation processes to ensure consistency

#### 4. Optimize for Query Patterns

- Design schemas around the most frequent query patterns
- Create appropriate indexes based on access patterns

#### 5. Use Composite Identifiers Consistently

- Maintain the same ID format across both databases
- Consider embedding database source in IDs (e.g., `pg-uuid` vs `neo-uuid`)

#### 6. Implement Monitoring for Both Databases

- Track query performance in both systems
- Monitor cross-database synchronization latency

#### 7. Design for Scale from the Start

- Consider sharding strategies for both PostgreSQL and Neo4j
- Implement read replicas for high-read workloads
- *Consider Neon or design postgresQL with extensibility for a post MVP sharding implementation*

#### 8. Leverage Database-Specific Features

- Use Neo4j's (or MemGraph) graph algorithms library for network analysis
- Utilize PostgreSQL's robust transactional capabilities for structured data

This comprehensive approach balances the strengths of both databases while maintaining data consistency and providing optimal performance for Penumbra's knowledge graph capabilities.

---

## Local KG Example (From my Dev Environment)

## Knowledge Graph Visualization

Markdown

graph TD

%% Define nodes with their types and key observations

LLM[("LLM (System)\n[TrainedOnTextData]")]

KG[("Knowledge Graph (System)\n[StructuredInfo]")]

NLU[("Natural Language\nUnderstanding (Capability)")]

SI[("Schema Inference\n(Capability)")]

%% Define relationships with descriptive labels

LLM -->|possesses| NLU

LLM -->|possesses| SI

LLM -->|maintains| KG

NLU -->|"enables creation of"| KG

SI -->|"enables flexible\nevolution of"| KG

%% Style definitions

classDef system fill:#f9f,stroke:#333,stroke-width:2px;

classDef capability fill:#bbf,stroke:#333,stroke-width:2px;

%% Apply styles

class LLM,KG system;

class NLU,SI capability;

%% Add key observations as notes

subgraph Observations

LLM\_obs["LLM Properties:

- Trained on vast text data
- Understands natural language
- Recognizes patterns
- Adapts to schemas"]

KG\_obs["Knowledge Graph Properties:

- Structured information
- Entities & relationships
- Queryable & traversable
- Semantic relationships"]

NLU\_obs["NLU Capabilities:

- Parse human language
- Extract entities & relationships
- Understand context"]

SI\_obs["Schema Inference:

- Work with data structures
- Infer relationships
- Flexible adaptation"]

end

%% Connect observations to nodes

LLM -.->LLM\_obs

KG -.->KG\_obs

```
NLU -. -> NLU_obs
SI -. -> SI_obs
```

## Diagram Explanation

This diagram represents a knowledge graph I created on my local dev environment, showing:

### 1. Node Types:

- **System nodes (pink):**
  - LLM and Knowledge Graph
- **Capability nodes (blue):**
  - Natural Language Understanding and Schema Inference

### 1. Relationships:

- Direct relationships shown with solid lines
- Observational properties shown with dotted lines to text boxes

### 1. Properties:

- Each node's key properties are listed in associated text boxes
- Properties help explain the role and capabilities of each component

### 1. Visual Hierarchy:

- Central nodes represent core systems
- Supporting capabilities branch out
- Observations provide additional context

---

▼ MCP Servers

---

# Reddit Researcher MCP

---

## 1 Test Current Version

---

### ☐ Test Current Version

- Follow the setup script for the Reddit app.
  - Access the provided GitHub link to the setup script
  - Run the script to create a Reddit app and generate credentials
  - Follow instructions for configuring the app in the repository
  - Verify permissions and security settings as outlined

---

▼ AI Staff

---

 Page

# Founder Mode Mentality

---

## 1 Assessment

### 1.1 Analysis

#### 1.1.1 AI staff infrastructure for Penumbra

### 1.2 Approach

#### 1.2.1 Don't just build AI agents

##### 1.2.1.1 1. 1.2.1.1 Cross-pollinate insights1.2.1.1

##### 1.2.1.2 2. 1.2.1.2 Self-improve without human intervention1.2.1.2

##### 1.2.1.3 3. 1.2.1.3 Skip hierarchies entirely1.2.1.3

##### 1.2.1.4 4. 1.2.1.4 Operate 24/7 in parallel1.2.1.4

### 1.3 Alignment

### 1.4 Reasoning

## 2 Manager Mode Mentality

### 2.1 Situation Assessment

### 2.2 Action Plan

### 2.3 Resource Allocation

### 2.4 Risk Management



## Assessment

### Analysis

#### AI staff infrastructure for Penumbra

I see it as essentially creating an "org chart" of AI agents to handle specialized tasks. This approach to operationalizing AI as team members rather than just tools is exactly the kind of unconventional thinking that drives innovation.

We're not just automating tasks;  
we're reimagining how teams function with AI as core collaborators.

### Approach

#### Don't just build AI agents

- Build an intelligent neural network for your company where these agents can:

##### 1. Cross-pollinate insights

- Have our [AI VP of Product](#) feed data to our [Metrics Analyst](#) who then informs our [Trend Analyst](#) in a continuous loop

##### 1. Self-improve without human intervention

- Build agents that can critique and improve each other's work

##### 1. Skip hierarchies entirely

- Let AI agents pull context directly from anywhere in your system, unlike human teams constrained by org charts

##### 1. Operate 24/7 in parallel

- Run multiple experiments simultaneously with different agents tackling different angles

Rather than following the conventional microservice approach, consider building a "hive mind" architecture where agents share a unified context layer but have specialized output functions.

### Alignment

This approach perfectly aligns with the vision for Penumbra as "the Brain for Agents" and "Intelligent Ops."

We're not just building *individual AI tools*;

We're creating a cognitive infrastructure that fundamentally transforms how companies operate.

This is the foundation for truly intelligent operations that can scale exponentially without the linear constraints of human teams.

## Reasoning

The standard approach would be to build isolated AI services that mimic human roles. But founder thinking means seeing what others don't

in this case, that AI agents can operate with fundamentally different organizational principles than humans.

By treating our AI staff as a collective intelligence rather than isolated specialists, we're creating something that couldn't exist in a human organization. This is the essence of founder mode thinking - not just digitizing existing processes but reimagining what's possible when freed from traditional constraints.

## Manager Mode Mentality

### Situation Assessment

We've developed a structured approach to implementing AI staff roles with defined inputs, outputs, and metrics. The current plan includes specialized roles like AI QA Analyst, Copywriter, and Metrics Analyst with specific tasks mapped to each role. WE've established governance metrics and a week-by-week implementation plan.

### Action Plan

1. Prioritize roles based on immediate business impact and technical feasibility
2. Implement the AI QA Analyst first as it can provide immediate value in test failure analysis
3. Develop standardized templates for all system prompts to ensure consistency (DSPy & Portkey)
4. Create documentation for how human team members should interact with AI staff
5. Establish clear escalation paths for when AI outputs need human intervention

### Resource Allocation

- Dedicate 1-2 engineers (*how does this quantify in the age of Agents?*) to build the microservice infrastructure
- Allocate time from product managers to define clear input/output specifications
- Set aside 20% of implementation time for testing and refinement

- Create a dedicated Slack channel for feedback collection

## Risk Management

- Potential over-reliance on AI outputs without proper validation
- Risk of prompt drift over time leading to inconsistent results
- Technical challenges with Linear webhook integration
- Team adoption barriers if the system is difficult to use

## Timeline and Milestones

- Week 1: Complete pilot role implementation (AI QA Analyst)
- Week 2: Measure initial acceptance rates and refine prompts
- Week 3: Roll out second AI staff role (likely AI Copywriter)
- Week 4: Implement dashboard for monitoring AI staff performance
- Week 5: Complete documentation and training for broader team adoption

## Performance Metrics

- Acceptance Rate: Target >80% of outputs usable without edits
- Cycle Time: Reduce by 50% compared to fully human process
- Usage Rate: >70% of eligible tasks routed through AI staff
- Cost Efficiency: Calculate ROI based on time saved vs. implementation costs

---

### ▼ Tech Stack

#### Ideation 1

---

 Page

## Technology Ideation 1

---

### 1 Exploration

#### 1.1 Possible Frameworks by Use Case

#### 1.2 Exploring newer frameworks:

### 2 Possible Agent Architecture

#### 2.1 Core Orchestration Frameworks

#### 2.2 Google's Agent Ecosystem

#### 2.3 Specialized Components

## 2.4 Full Stack

### 2.4.1 Infrastructure & DevOps

### 2.4.2 Development Environment

### 2.4.3 AI/ML Infrastructure

### 2.4.4 Security, Validation & Compliance

### 2.4.5 Unit Testing

### 2.4.6 Testing Beyond Unit Tests

### 2.4.7 API Management & LLM Orchestration

### 2.4.8 Vector Databases & Storage

### 2.4.9 Data Pipeline

### 2.4.10 Frontend Enhancements

### 2.4.11 Agent-Specific Components

### 2.4.12 Documentation

---

## Exploration

### Possible Frameworks by Use Case

1. LangChain/LangGraph for agent orchestration
2. AutoGen framework for multi-agent systems
3. CrewAI for collaborative AI agents
4. DSPy for programmatic prompt engineering
5. Semantic Kernel for agent integration
6. Fixie.ai for agent infrastructure

### Exploring newer frameworks:

- LlamaIndex for knowledge integration
- Haystack for agent pipelines
- Flowise for visual agent building
- Langflow for flow-based agent design

## Possible Agent Architecture

### Core Orchestration Frameworks

- **LangChain/LangGraph:** For creating the fundamental agent workflows and state management
- **Flow-based tools** (like Flowise or Langflow): For visual development of agent pipelines

## Google's Agent Ecosystem

- **Google ADK** (Agent Development Kit): Google's new toolkit for building AI agents
- **A2A Protocol** (Agent-to-Agent): The communication protocol Google recently announced for standardizing how agents interact
- **AG2**: Likely referring to AgentGPT or a similar agent framework

## Specialized Components

- **Magnetic One**: For agent orchestration with stronger reasoning capabilities
- **Browser Integration**:
  - **O-Man (My Unreleased Local Agent)**: For browser automation capabilities
  - **Proxy Lite**: For web scraping and interaction capabilities

This multi-framework approach makes sense for Penumbra as "the Brain for Agents"

- we're essentially creating a meta-framework that can leverage the strengths of each tool while providing a unified interface for agent management.
  - The A2A protocol is particularly interesting since it's gaining industry support (over 50 partners including Salesforce and ServiceNow) and focuses on "opaque execution" where agents don't need to share internal states - just inputs and outputs.
- 

## Full Stack

### Infrastructure & DevOps

- **Kubernetes** alongside Docker Swarm for more complex orchestration needs
- **Lefthook** for git hooks
- **Redis** for caching/messaging
- **iKafka** for streaming
- **Terraform/Pulumi** for infrastructure as code
- **Observability stack**: Prometheus, Grafana, OpenTelemetry for monitoring Basic Agent Architectural performance
- **Vector DB**: Pinecone, Weaviate, or Milvus for embedding storage
- **Message Queue**: Kafka or RabbitMQ for async agent communication

### Development Environment

- **Dev Containers** for consistent environments
- **Nx** as an alternative to Turborepo worth considering
- **Husky** to complement lefthook for git hooks

### AI/ML Infrastructure

- **Model serving:** vLLM, Text Generation Inference, or Ray Serve
- **Ray** for distributed computing
- **Feature store:** for tracking agent performance features
- **Experiment tracking:** MLflow or Weights & Biases
- **Prompt versioning:** LangSmith or similar
- **W&B** (Weights & Biases) for experiment tracking
- **LangWatch** for LLM monitoring
- **Langfuse** for LLM observability
- **Opik** for performance optimization
- **PromptFoo** for prompt testing

## Security, Validation & Compliance

- **Auth system:** Clerk (alternatives: Auth0, or Supabase Auth)
- **Rate limiting** for API endpoints
- **Input validation/sanitization** frameworks
- **Audit logging** for agent actions
- **Zod** for TypeScript validation
- **Pydantic** for Python validation

## Unit Testing

- **Jest** for JavaScript testing
- **Pytest** for Python testing

## Testing Beyond Unit Tests

- **Integration tests** for agent workflows
- **Prompt testing** framework
- **Agent simulation environment** for testing multi-agent interactions
- **Chaos testing** for resilience

## API Management & LLM Orchestration

- **Portkey** for API management and routing (instead of Diamond)
- **OpenRouter** for model access
- **LiteLLM** for LLM provider abstraction

## Vector Databases & Storage

- **PGVector** for PostgreSQL vector extensions
- **Pinecone** for managed vector search
- **Weaviate** for semantic search

- **Turso** for edge SQLite
- **Neon** for serverless Postgres
- **AstraDB** for scalable NoSQL
- **MongoDB** for document storage

## Data Pipeline

- **ETL/ELT tools:** Airbyte, dbt
- **Streaming data:** Kafka Streams, Flink
- **Crawl4AI** for web crawling
- **LightRAG** for efficient retrieval
- **Data lineage tracking**

## Frontend Enhancements

- **Storybook** for component development
- **Supernova** for design system management
- **Shadcn** for UI components
- **21stDev** for developer experience
- **Vitest** for faster testing
- **Tanstack Query** for data fetching
- **Zustand/Jotai** for state management

## Agent-Specific Components

- **Agent memory store:** structured storage for agent context
- **Tool registry:** for managing agent capabilities
- **Prompt template management system**
- **Agent debugging console/UI**

## Documentation

- **OpenAPI** for API documentation
- **Docusaurus** for technical documentation
- **Storybook** for component documentation

---

## Stack Proposal 1

---

# Penumbra: The Brain for Agents - Technical Architecture

---

- 1 Overview
- 2 Table of Contents
- 3 Frontend Applications
- 4 Frontend Technology
  - 4.1 UI Components & Design
  - 4.2 Animation & Visualization
  - 4.3 Development Tools
  - 4.4 State Management & Data Fetching
- 5 API & Backend Services
- 6 AI Agent Infrastructure
  - 6.1 LLM Orchestration
  - 6.2 Agent Frameworks
- 7 Data Infrastructure
  - 7.1 Vector Stores
  - 7.2 Databases
  - 7.3 Data Processing
  - 7.4 Messaging & Cache
- 8 DevOps & Infrastructure
  - 8.1 Build & Package Management
  - 8.2 Containerization & Deployment
  - 8.3 CI/CD
  - 8.4 Monitoring & Observability
- 9 Security & Validation
- 10 Documentation
- 11 Implementation Roadmap
  - 11.1 Phase 1: Foundation
  - 11.2 Phase 2: Backend Infrastructure
  - 11.3 Phase 3: AI Infrastructure
  - 11.4 Phase 4: Advanced Features



# Overview

Penumbra is a comprehensive platform for building, orchestrating, and managing AI agents. This document outlines the complete technical architecture and tooling choices for the platform.

## Table of Contents

1. Frontend Applications
2. Frontend Technology
3. API & Backend Services
4. AI Agent Infrastructure
5. Data Infrastructure
6. DevOps & Infrastructure
7. Security & Validation
8. Documentation
9. Implementation Roadmap

## Frontend Applications

- **Next.js WebUI** - Primary user interface for interacting with Penumbra
- **Agent Dashboard** - Management console for monitoring and configuring AI agents
- **Prompt Studio** - Interactive environment for designing and testing prompts
- **Analytics & Monitoring** - Visualization of agent performance and system metrics

## Frontend Technology

### UI Components & Design

- **Radix UI** - Unstyled, accessible UI primitives
- **ShadCN** - Component library built on Radix
- **Tailwind 4** - Utility-first CSS framework
- **21st Dev** - Developer experience enhancements
- **Supernova** - Design system management
- **Composio** - Visual composition tool

### Animation & Visualization

- **Framer Motion** - Animation library

- **Anime.js** - JavaScript animation engine
- **Three.js / R3F / Drei** - 3D visualization libraries

## Development Tools

- **Storybook** - Component development and documentation

## State Management & Data Fetching

- **Zustand/Jotai** - State management libraries
- **Tanstack Query** - Data fetching and caching

## API & Backend Services

- **FastAPI** - Python-based API framework for agent services
- **NestJS** - Node.js framework for real-time services
- **GraphQL** - Query language for APIs
- **API Gateway** - Unified entry point for all services

## AI Agent Infrastructure

### LLM Orchestration

- **Portkey** - API management and routing
- **OpenRouter** - Model access and routing
- **LiteLLM** - LLM provider abstraction
- **LangWatch** - LLM monitoring
- **Langfuse** - Observability for LLMs
- **PromptFoo** - Prompt testing framework

### Agent Frameworks

- **LangChain** - Framework for LLM applications
- **LangGraph** - Graph-based framework for agent workflows
- **CrewAI** - Multi-agent orchestration
- **AutoGen** - Conversational agent framework
- **Semantic Kernel** - Microsoft's AI orchestration framework
- **Google ADK** - Google's Agent Development Kit
- **A2A Protocol** - Agent-to-Agent communication protocol
- **Flowise/Langflow** - Visual agent builders
- **O-Man/ProxyLite** - Browser automation and web interaction

## Data Infrastructure

## Vector Stores

- **Pinecone** - Managed vector database
- **Weaviate** - Semantic vector search
- **PGVector** - PostgreSQL vector extensions

## Databases

- **PostgreSQL/Neon** - Relational database (serverless option)
- **MongoDB** - Document database
- **Turso** - Edge SQLite database
- **AstraDB** - Scalable NoSQL database

## Data Processing

- **dbt** - Data transformation
- **Ray** - Distributed computing framework
- **LightRAG** - Efficient retrieval augmented generation
- **Crawl4AI** - Web crawling for AI

## Messaging & Cache

- **Redis** - In-memory data store and messaging
- **iKafka** - Event streaming platform

## DevOps & Infrastructure

### Build & Package Management

- **Turborepo** - Monorepo build system
- **PNPM** - Fast, disk space efficient package manager
- **PDM** - Python package manager
- **uv Python** - Fast Python package installer and resolver

### Containerization & Deployment

- **Docker** - Containerization platform
- **Swarm** - Container orchestration
- **TSDProxy** - TypeScript-first proxy
- **NGinx** - Web server and reverse proxy
- **Dev Containers** - Consistent development environments

### CI/CD

- **GitHub Actions** - Continuous integration and deployment

- **Commitizen** - Standardized commit messages
- **GPTLint** - AI-powered linting
- **Lint-staged** - Run linters on staged files
- **Lefthook** - Git hooks manager

## Monitoring & Observability

- **Prometheus** - Monitoring system
- **Grafana** - Observability platform
- **W&B** (Weights & Biases) - ML experiment tracking
- **Opik** - Performance optimization

## Security & Validation

- **Clerk** - Authentication and user management
- **Zod** - TypeScript-first schema validation
- **Pydantic** - Data validation for Python

## Documentation

- **Mintlify** - Documentation platform
- **Storybook** - Component documentation

## Implementation Roadmap

### Phase 1: Foundation

1. Set up monorepo structure with Turborepo, PNPM, and PDM
2. Establish CI/CD pipelines with GitHub Actions
3. Configure Docker and development containers
4. Implement core Next.js application with Tailwind and Radix UI

### Phase 2: Backend Infrastructure

1. Develop FastAPI and NestJS services
2. Set up database infrastructure (PostgreSQL, MongoDB, Redis)
3. Implement authentication with Clerk
4. Configure monitoring with Prometheus and Grafana

### Phase 3: AI Infrastructure

1. Integrate LLM orchestration (Portkey, LiteLLM)
2. Set up vector databases (Pinecone, PGVector)

3. Implement core agent frameworks (LangChain, LangGraph)
4. Develop RAG capabilities with LightRAG

## Phase 4: Advanced Features

1. Multi-agent orchestration with CrewAI and AutoGen
2. Browser automation with O-Man/ProxyLite
3. Advanced analytics and monitoring
4. 3D visualization with Three.js/R3F

## Phase 5: Optimization & Scale

1. Performance optimization with Opik
2. Distributed computing with Ray
3. Enhanced observability with Langfuse and W&B
4. Production hardening and security enhancements

---

*Penumbra: Intelligent Ops for the AI-native enterprise*

---

## Stack Proposal Diagram 1

---

 Page

## Stack Proposal Diagram 1

---

---

Markdown

## flowchart TB

%% Main Platform

Platform["PENUMBRA PLATFORM"]

%% Main Sections

Frontend["FRONTEND APPLICATIONS"]

FrontendTech["FRONTEND TECHNOLOGY"]

Backend["API & BACKEND SERVICES"]

AIInfra["AI AGENT INFRASTRUCTURE"]

DataInfra["DATA INFRASTRUCTURE"]

DevOps["DEVOPS & INFRASTRUCTURE"]

Security["SECURITY & VALIDATION"]

Docs["DOCUMENTATION"]

%% Frontend Applications

FE1["Next.js WebUI"]

FE2["Agent Dashboard"]

FE3["Prompt Studio"]

FE4["Analytics & Monitoring"]

%% Frontend Technology Groups

FT1["UI Components & Design"]

FT2["Animation & Visualization"]

FT3["State & Data Fetching"]

%% Frontend Tech Components

FT1\_1["Radix UI / ShadCN"]

FT1\_2["Tailwind 4 / 21st Dev"]

FT1\_3["Supernova / Composio / Storybook"]

FT2\_1["Framer Motion / Anime.js"]

FT2\_2["Three.js / R3F / Drei"]

FT3\_1["Zustand / Jotai"]

FT3\_2["Tanstack Query"]

%% Backend Services

BE1["FastAPI"]

BE2["NestJS"]

BE3["GraphQL"]

BE4["API Gateway"]

%% AI Infrastructure

AI1["LLM ORCHESTRATION"]

AI2["AGENT FRAMEWORKS"]

%% LLM Orchestration

AI1\_1["Portkey / OpenRouter"]

AI1\_2["LiteLLM / LangWatch"]

AI1\_3["Langfuse / PromptFoo"]

%% Agent Frameworks

AI2\_1["LangChain / LangGraph"]

```
AI2_2["CrewAI / AutoGen"]
AI2_3["Semantic Kernel"]
AI2_4["Google ADK / A2A Protocol"]
AI2_5["Flowise / Langflow"]
AI2_6["O-Man / ProxyLite"]
```

```
%% Data Infrastructure
Data1["VECTOR STORES"]
Data2["DATABASES"]
Data3["DATA PROCESSING"]
Data4["MESSAGING & CACHE"]
```

```
%% Vector Stores
Data1_1["Pinecone"]
Data1_2["Weaviate"]
Data1_3["PGVector"]
```

```
%% Databases
Data2_1["PostgreSQL / Neon"]
Data2_2["MongoDB"]
Data2_3["Turso"]
Data2_4["AstraDB"]
```

```
%% Data Processing
Data3_1["dbt"]
Data3_2["Ray"]
Data3_3["LightRAG"]
Data3_4["Crawl4AI"]
```

```
%% Messaging & Cache
Data4_1["Redis"]
Data4_2["iKafka"]
```

```
%% DevOps
Dev1["BUILD & PACKAGE"]
Dev2["CONTAINERIZATION"]
Dev3["CI/CD"]
Dev4["MONITORING"]
```

```
%% Build & Package
Dev1_1["Turborepo / PNPM"]
Dev1_2["PDM / uv Python"]
```

```
%% Containerization
Dev2_1["Docker / Swarm"]
Dev2_2["TSDProxy / NGnix"]
Dev2_3["Dev Containers"]
```

```
%% CI/CD
Dev3_1["GitHub Actions"]
Dev3_2["Commitizen / GPTLint"]
```



```
Dev3_3["Lint-staged / Lefthook"]
```

```
%% Monitoring
```

```
Dev4_1["Prometheus / Grafana"]
```

```
Dev4_2["W&B / Opik"]
```

```
%% Security & Validation
```

```
Sec1["Clerk"]
```

```
Sec2["Zod / Pydantic"]
```

```
%% Documentation
```

```
Doc1["Mintlify"]
```

```
Doc2["Storybook"]
```

```
%% Main Flow
```

```
Platform --> Frontend
```

```
Platform --> FrontendTech
```

```
Platform --> Backend
```

```
Platform --> AIInfra
```

```
Platform --> DataInfra
```

```
Platform --> DevOps
```

```
Platform --> Security
```

```
Platform --> Docs
```

```
%% Frontend Applications
```

```
Frontend --- FE1
```

```
Frontend --- FE2
```

```
Frontend --- FE3
```

```
Frontend --- FE4
```

```
%% Frontend Technology
```

```
FrontendTech --- FT1
```

```
FrontendTech --- FT2
```

```
FrontendTech --- FT3
```

```
FT1 --- FT1_1
```

```
FT1 --- FT1_2
```

```
FT1 --- FT1_3
```

```
FT2 --- FT2_1
```

```
FT2 --- FT2_2
```

```
FT3 --- FT3_1
```

```
FT3 --- FT3_2
```

```
%% Backend Services
```

```
Backend --- BE1
```

```
Backend --- BE2
```

```
Backend --- BE3
```

```
Backend --- BE4
```

```
%% AI Infrastructure
```

```
AIInfra --- AI1
```

AIInfra --- AI2

AI1 --- AI1\_1

AI1 --- AI1\_2

AI1 --- AI1\_3

AI2 --- AI2\_1

AI2 --- AI2\_2

AI2 --- AI2\_3

AI2 --- AI2\_4

AI2 --- AI2\_5

AI2 --- AI2\_6

%% Data Infrastructure

DataInfra --- Data1

DataInfra --- Data2

DataInfra --- Data3

DataInfra --- Data4

Data1 --- Data1\_1

Data1 --- Data1\_2

Data1 --- Data1\_3

Data2 --- Data2\_1

Data2 --- Data2\_2

Data2 --- Data2\_3

Data2 --- Data2\_4

Data3 --- Data3\_1

Data3 --- Data3\_2

Data3 --- Data3\_3

Data3 --- Data3\_4

Data4 --- Data4\_1

Data4 --- Data4\_2

%% DevOps

DevOps --- Dev1

DevOps --- Dev2

DevOps --- Dev3

DevOps --- Dev4

Dev1 --- Dev1\_1

Dev1 --- Dev1\_2

Dev2 --- Dev2\_1

Dev2 --- Dev2\_2

Dev2 --- Dev2\_3

Dev3 --- Dev3\_1

Dev3 --- Dev3\_2

```
Dev3 --- Dev3_3

Dev4 --- Dev4_1
Dev4 --- Dev4_2

%% Security & Validation
Security --- Sec1
Security --- Sec2

%% Documentation
Docs --- Doc1
Docs --- Doc2

%% Styling
classDef platform fill:#f9f,stroke:#333,stroke-width:2px
classDef section fill:#bbf,stroke:#333,stroke-width:1px
classDef subsection fill:#ddf,stroke:#333,stroke-width:1px
classDef component fill:#fff,stroke:#333,stroke-width:1px

class Platform platform
class
Frontend,FrontendTech,Backend,AIInfra,DataInfra,DevOps,Security,Docs
section
class
FE1,FE2,FE3,FE4,FT1,FT2,FT3,BE1,BE2,BE3,BE4,AI1,AI2,Data1,Data2,Data3,Data4,Dev1,Dev2,Dev3,Dev4,Sec1,Sec2,Doc1,Doc2 subsection
class
FT1_1,FT1_2,FT1_3,FT2_1,FT2_2,FT3_1,FT3_2,AI1_1,AI1_2,AI1_3,AI2_1,AI2_2,AI2_3,AI2_4,AI2_5,AI2_6,Data1_1,Data1_2,Data1_3,Data2_1,Data2_2,Data2_3,Data2_4,Data3_1,Data3_2,Data3_3,Data3_4,Data4_1,Data4_2,Dev1_1,Dev1_2,Dev2_1,Dev2_2,Dev2_3,Dev3_1,Dev3_2,Dev3_3,Dev4_1,Dev4_2 component
```

## ▼ DevSecOps

 Page

# GitHub Repository Security Best Practices

1 InfoSec & DevSecOps for GitHub Repos

2 Access Control & Repository Configuration

2.1 1. 2.1 Implement Branch Protection Rules

2.2 2. 2.2 Enable Two-Factor Authentication (2FA)

2.3 3. 2.3 Use Fine-Grained Access Control

#### 2.4 4. 2.4 Set Up Security Policies

### 3 Code Security

#### 3.1 5. 3.1 Implement Code Scanning

#### 3.2 6. 3.2 Secret Management

#### 3.3 7. 3.3 Dependency Management

#### 3.4 8. 3.4 Secure Coding Practices

### 4 CI/CD Pipeline Security

#### 4.1 9. 4.1 Secure GitHub Actions

#### 4.2 10. 4.2 Implement Security Testing in CI/CD

#### 4.3 11. 4.3 Deploy with Least Privilege

### 5 Compliance & Documentation

#### 5.1 12. 5.1 License Compliance

#### 5.2 13. 5.2 Security Documentation

#### 5.3 14. 5.3 Compliance Standards

#### 5.4 Monitoring & Incident Response

#### 5.5 15. 5.5 Security Monitoring

#### 5.6 16. 5.6 Vulnerability Management

#### 5.7 17. 5.7 Incident Response Plan

### 6 Supply Chain Security

#### 6.1 18. 6.1 Software Bill of Materials ( [6.1 SBOM](#) 6.1 )

#### 6.2 19. 6.2 Verify Dependencies

#### 6.3 20. 6.3 Secure Build Processes

### 7 Shift-Left Security Approach

#### 7.1 21. 7.1 Early Security Integration

#### 7.2 22. 7.2 Security as Code

---

## InfoSec & DevSecOps for GitHub Repos

These standards will help secure projects we are developing

## Access Control & Repository Configuration

### 1. Implement Branch Protection Rules

- Require pull request reviews before merging
- Require status checks to pass before merging
- Require signed commits

- Restrict who can push to matching branches

## 2. Enable Two-Factor Authentication (2FA)

- Require 2FA for all organization members
- Consider hardware security keys for critical repositories

## 3. Use Fine-Grained Access Control

- Limit admin access to essential personnel only
- Apply the principle of least privilege
- Regularly audit access permissions

## 4. Set Up Security Policies

- Create a [SECURITY.md](#) file with vulnerability reporting guidelines
- Define a clear process for security issue disclosure

# Code Security

## 1. Implement Code Scanning

- Enable GitHub Advanced Security (if available)
- Configure CodeQL analysis for automated vulnerability detection
- Set up custom code scanning alerts

## 2. Secret Management

- Use GitHub Secrets for storing sensitive information
- Implement pre-commit hooks to prevent secrets from being committed
- Use tools like GitGuardian or Gitleaks to scan for exposed secrets

## 3. Dependency Management

- Enable Dependabot alerts and security updates
- Configure dependency review for pull requests
- Regularly update dependencies

## 4. Secure Coding Practices

- Establish coding standards that emphasize security
  - Document secure patterns for common operations
  - Create security-focused code review checklists
-

# CI/CD Pipeline Security

## 1. Secure GitHub Actions

- Pin actions to specific SHA hashes instead of tags
- Use trusted actions from the marketplace or create your own
- Limit permissions for GitHub Actions workflows

## 2. Implement Security Testing in CI/CD

- **SAST** (Static Application Security Testing)
- **DAST** (Dynamic Application Security Testing)
- **SCA** (Software Composition Analysis)
- Container scanning for Docker images

## 3. Deploy with Least Privilege

- Use dedicated deployment credentials
  - Rotate deployment keys regularly
  - Implement time-limited access tokens
- 

# Compliance & Documentation

## 1. License Compliance

- Include appropriate ``LICENSE`` files
- Ensure third-party dependencies comply with your license
- Document license requirements

## 2. Security Documentation

- Maintain up-to-date security documentation
- Document security controls and their implementations
- Create incident response procedures

## 3. Compliance Standards

- Implement controls for relevant standards ([GDPR](#), [SOC 2](#), etc.)
  - Document compliance measures
-

# Monitoring & Incident Response

## 1. Security Monitoring

- Set up alerts for suspicious activities
- Monitor for unusual commit patterns
- Track and analyze security events

## 2. Vulnerability Management

- Establish a process for addressing security vulnerabilities
- Define severity levels and response times
- Track remediation progress

## 3. Incident Response Plan

- Document steps to take when security incidents occur
  - Define roles and responsibilities
  - Practice incident response scenarios
- 

# Supply Chain Security

## 1. Software Bill of Materials (SBOM)

- Generate and maintain SBOMs for all projects
- Use tools like CycloneDX or SPDX

## 2. Verify Dependencies

- Use lockfiles to pin dependency versions
- Verify package integrity with checksums
- Consider using private package registries for critical dependencies

## 3. Secure Build Processes

- Implement reproducible builds
- Sign build artifacts
- Document build environment requirements

# Shift-Left Security Approach

I propos "**Shift Left**" in our **SOPs**, which is an excellent security approach:

## 1. Early Security Integration

- Integrate security checks at the earliest stages of development
- Provide security training for developers / contributors
- Create security champions within development teams

## 2. Security as Code

- Define security policies as code
- Automate security checks
- Version control security configurations

These best practices will help establish a robust security posture for our GitHub repositories while maintaining development velocity.

---

---

 Page

# Repository Configuration Files for Security & Best Practices

---

## 1 Git Configuration Files

### 1.1 .gitignore

## 2 Environment & Dependency Management

### 2.1 .env 2.1 (Never commit this file)

### 2.2 .npmrc

### 2.3 .nvmrc

## 3 Code Quality & Security Tools

### 3.1 .prettierrc

### 3.2 .eslintrc

## 4 GitHub Workflows & Actions

### 4.1 GitHub Actions Workflow for Security Scanning

### 4.2 MegaLinter Configuration4.2

### 4.3 WhiteSource Bolt Configuration



## 5 Additional Security Files

5.1 `dependabot.yml`

5.2 `SECURITY.md`

5.3 `.dockerignore`

5.4 `.pre-commit-config.yaml`

---

## Git Configuration Files

### `.gitignore`

- Exclude sensitive files and directories ( `.env` , credentials, logs)
- Exclude build artifacts, dependencies, and cache directories
- Include language/framework-specific exclusions (node\_modules, **pycache**, etc.)
- Example for Node.js/React projects:

Text

```
# Environment variables
.env
.env.local
.env.development.local
.env.test.local
.env.production.local

# Dependencies
/node_modules
/.pnp
.pnp.js

# Build artifacts
/build
/dist
/.next/
/out/

# Cache
.cache/
.npm
.eslintcache

# Logs
logs
*.log
npm-debug.log*
yarn-debug.log*
yarn-error.log*
pnpm-debug.log*
pnpm-error.log*

# IDE specific files
.idea/
.vscode/
*.swp
*.sw0
```

.gitattributes

- Enforce consistent line endings
  - Define binary files to prevent merge conflicts
  - Set merge strategies for specific file types
- plaintext

Text

```
# Set default behavior to automatically normalize line endings
* text=auto

# Explicitly declare text files to be normalized
*.js text
*.jsx text
*.ts text
*.tsx text
*.json text
*.md text
*.css text
*.html text

# Denote binary files that should not be modified
*.png binary
*.jpg binary
*.gif binary
*.ico binary
*.woff binary
*.woff2 binary

# Apply specific merge strategies
package-lock.json merge=ours
yarn.lock merge=ours
```

## Environment & Dependency Management

### **.env** (Never commit this file)

- Store environment-specific secrets and configuration
- Create **.env.example** with placeholders to commit instead

Text

```
# .env.example (safe to commit)
DATABASE_URL=postgres://username:password@localhost:5432/database
API_KEY=your_api_key_here
JWT_SECRET=your_jwt_secret_here
```

### **.npmrc**

- Configure npm to use secure package sources
- Set up authentication for private registries
- Enable security audits

Text

```
registry=https://registry.npmjs.org/  
audit=true  
audit-level=high  
save-exact=true  
fund=false  
engine-strict=true
```

### .nvmrc

- Lock Node.js version for consistent development environments
- Prevents security issues from version mismatches

Text

22.12.12

## Code Quality & Security Tools

### .prettierrc

- Enforce consistent code formatting
- Prevent formatting-related security issues

JSON

```
{  
  "singleQuote": true,  
  "trailingComma": "es5",  
  "printWidth": 100,  
  "tabWidth": 2,  
  "semi": true,  
  "arrowParens": "avoid"  
}
```

### .eslintrc

- Enforce secure coding practices
- Detect potential security vulnerabilities
- Include security-focused plugins

JSON

```
{
  "extends": [
    "eslint:recommended",
    "plugin:@typescript-eslint/recommended",
    "plugin:react/recommended",
    "plugin:react-hooks/recommended",
    "plugin:security/recommended"
  ],
  "plugins": [
    "@typescript-eslint",
    "react",
    "react-hooks",
    "security"
  ],
  "rules": {
    "security/detect-object-injection": "error",
    "security/detect-non-literal-regexp": "error",
    "security/detect-unsafe-regex": "error",
    "security/detect-buffer-noassert": "error",
    "security/detect-eval-with-expression": "error",
    "security/detect-no-csrf-before-method-override": "error",
    "security/detect-possible-timing-attacks": "error"
  }
}
```

## GitHub Workflows & Actions

### GitHub Actions Workflow for Security Scanning

YAML

```
name: Security Scan

on:
  push:
    branches: [main, develop]
  pull_request:
    branches: [main, develop]
  schedule:
    - cron: '0 0 * * 0' # Weekly scan

jobs:
  security-scan:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v3
        with:
          fetch-depth: 0

      - name: CodeQL Analysis
        uses: github/codeql-action/init@v2
        with:
          languages: javascript, typescript

      - name: Perform CodeQL Analysis
        uses: github/codeql-action/analyze@v2

      - name: Dependency Review
        uses: actions/dependency-review-action@v3

      - name: Run npm audit
        run: npm audit --audit-level=high

      - name: SAST with SonarCloud
        uses: SonarSource/sonarcloud-github-action@v1.9
        env:
          GITHUB_TOKEN: ${ secrets.GITHUB_TOKEN }
          SONAR_TOKEN: ${ secrets.SONAR_TOKEN }
```

## MegaLinter Configuration

```

YAML
name: MegaLinter

on:
  push:
    branches: [main, develop]
  pull_request:
    branches: [main, develop]

jobs:
  megalint:
    name: MegaLinter
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v3
        with:
          fetch-depth: 0

      - name: MegaLinter
        uses: megalinter/megalinter/flavors/javascript@v6
        env:
          GITHUB_TOKEN: ${ secrets.GITHUB_TOKEN }
          VALIDATE_ALL_CODEBASE: ${ github.event_name == 'push' &&
github.ref == 'refs/heads/main' }}
          JAVASCRIPT_ES_CONFIG_FILE: .eslintrc.json
          MARKDOWN_MARKDOWNLINT_CONFIG_FILE: .markdownlint.json
          YAML_YAMLLINT_CONFIG_FILE: .yamllint.yml
          FILTER_REGEX_EXCLUDE: '(\.git|node_modules|\.vscode)'

```

## WhiteSource Bolt Configuration

YAML

```
name: WhiteSource Scan

on:
  push:
    branches: [main]
  pull_request:
    branches: [main]
  schedule:
    - cron: '0 0 * * 0' # Weekly scan

jobs:
  whitesource:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v3

      - name: WhiteSource Unified Agent Scan
        uses: whitesource/actions@v1
        with:
          wssURL: ${ secrets.WSS_URL }
          apiKey: ${ secrets.WSS_API_KEY }
          productName: 'My Product'
          projectName: ${ github.repository }
          includes: '**/*.js **/*.jsx **/*.ts **/*.tsx'
          excludes: '**/node_modules/** **/dist/**'
```

## Additional Security Files

[dependabot.yml](#)



#### YAML

```
version: 2
updates:
  - package-ecosystem: "npm"
    directory: "/"
    schedule:
      interval: "weekly"
    open-pull-requests-limit: 10
    versioning-strategy: increase
    commit-message:
      prefix: "deps"
      include: "scope"
    labels:
      - "dependencies"
      - "security"

  - package-ecosystem: "github-actions"
    directory: "/"
    schedule:
      interval: "weekly"
    labels:
      - "ci-cd"
      - "security"
```

## SECURITY.md

#### Markdown

### # Security Policy

#### ## Supported Versions

Version	Supported
1.0.x	:white_check_mark:
< 1.0	:x:

#### ## Reporting a Vulnerability

Please report security vulnerabilities to [security@yourcompany.com](mailto:security@yourcompany.com).

Do not report security vulnerabilities through public GitHub issues.

You will receive a response within 48 hours. If the vulnerability is accepted, we'll work on a fix and release it according to our security release process.

## .dockerignore

Text

```
# Git
.git
.gitignore
.github
# Node.js
node_modules
npm-debug.log
yarn-error.log
# Environment
.env
.env.*
# Development
*.md
LICENSE
.dockerignore
Dockerfile
docker-compose*
.eslintrc
.prettierrc
.vscode
```

[.pre-commit-config.yaml](#)

YAML

```
repos:
  - repo: https://github.com/pre-commit/pre-commit-hooks
    rev: v4.4.0
    hooks:
      - id: trailing-whitespace
      - id: end-of-file-fixer
      - id: check-yaml
      - id: check-added-large-files
      - id: check-json
      - id: detect-private-key
      - id: no-commit-to-branch
        args: [--branch, main]

  - repo: https://github.com/gitleaks/gitleaks
    rev: v8.16.3
    hooks:
      - id: gitleaks

  - repo: https://github.com/eslint/eslint
    rev: v8.41.0
    hooks:
      - id: eslint
        files: \.(js|ts|jsx|tsx)$
        types: [file]
        additional_dependencies:
          - eslint
          - eslint-plugin-security
```

---

 Page

# Implementing dotenv-vault and Tailscale for Repository Security

---

## 1 Setup guide1

### 1.1 dotenv-vault Implementation

#### 1.2 1. Initial Setup

#### 1.3 2. Configuration Files

##### 1.3.1 .env.vault

##### 1.3.2 .env

##### 1.3.3 .env.example

- 1.4 3. GitHub Workflow Integration
  - 1.5 4. Package.json Scripts
  - 1.6 5. dotenv-vault Commands for Team Usage
  - 1.7 Working with environments
  - 2 Tailscale Implementation
    - 2.1 1. GitHub Actions Integration for Tailscale
    - 2.2 2. Docker Compose with Tailscale
    - 2.3 3. Tailscale Configuration File
    - 2.4 4. GitHub Repository Settings
    - 2.5 5. Combined Implementation for CI/CD
    - 2.6 6. Security Best Practices
- 

## Setup guide

### dotenv-vault Implementation

**dotenv-vault** is a secure environment variable manager that allows you to encrypt and share your environment variables across development teams



GitHub - dotenv-org/dotenv-vault: sync .env files—from the creator of `dotenv`.

<https://github.com/dotenv-org/dotenv-vault>

### 1. Initial Setup

```
Bash
# Install dotenv-vault CLI
npm install -g dotenv-vault

# Initialize in your project
cd your-project
dotenv-vault new
```

### 2. Configuration Files

#### `.env.vault`

This file will be created by dotenv-vault and should be committed to your repository. It contains encrypted environment variables.

#### `.env`

Your local environment file (should be in your `.gitignore`):

Text

```
# Development environment variables
DATABASE_URL=postgres://localhost:5432/mydb
API_KEY=local_development_key
```

### `.env.example`

Commit this to show required variables:

Text

```
# Required environment variables (without actual values)
DATABASE_URL=
API_KEY=
```

## 3. GitHub Workflow Integration

YAML

```
# .github/workflows/deploy.yml
name: Deploy with dotenv-vault
on:
  push:
    branches: [main]
jobs:
  deploy:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3

      - name: Setup Node.js
        uses: actions/setup-node@v3
        with:
          node-version: '18'

      - name: Install dependencies
        run: npm ci

      - name: Load environment variables
        run: npx dotenv-vault@latest pull production
        env:
          DOTENV_KEY: ${ secrets.DOTENV_KEY }

      - name: Build and deploy
        run: npm run build && npm run deploy
```

## 4. Package.json Scripts

JSON

```
"scripts": {  
  "dev": "dotenv -e .env next dev",  
  "build": "dotenv -e .env.vault next build",  
  "start": "dotenv -e .env.vault next start"  
}
```

## 5. dotenv-vault Commands for Team Usage

I will create runbooks so team members can press play button in the [README.md](#) file and automate setup

Add these to your README.md:

Markdown

### ## Environment Variables

We use dotenv-vault for secure environment management.

#### ### Setup

```
```bash  
# Install dotenv-vault CLI  
npm install -g dotenv-vault  
# Login to your team's vault  
dotenv-vault login  
# Pull the latest environment variables  
dotenv-vault pull
```

## Working with environments

Bash

```
# Create a new environment  
dotenv-vault open production  
  
# Push local changes to the vault  
dotenv-vault push  
  
# Build the encrypted .env.vault file for CI/CD  
dotenv-vault build
```

## Tailscale Implementation

Tailscale provides secure networking for your infrastructure, allowing secure access to services without exposing them to the public internet.



Tailscale · Best VPN Service for Secure Networks

<https://tailscale.com/>

## 1. GitHub Actions Integration for Tailscale

```

YAML
# .github/workflows/tailscale.yml
name: Deploy with Tailscale Access

on:
  push:
    branches: [main]

jobs:
  deploy:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3

      - name: Setup Tailscale
        uses: tailscale/github-action@v2
        with:
          oauth-client-id: ${ secrets.TS_OAUTH_CLIENT_ID }
          oauth-secret: ${ secrets.TS_OAUTH_SECRET }
          tags: tag:ci

      - name: Deploy to internal infrastructure
        run: |
          # Now you can access internal resources via Tailscale
          curl https://internal-service.your-tailnet.ts.net/deploy \
            -H "Authorization: Bearer ${ secrets.DEPLOY_TOKEN }" \
            --data-binary @deployment.json

```

## 2. Docker Compose with Tailscale

```

YAML
# docker-compose.yml
version: '3'
services:
  app:
    build: .
    environment:
      - NODE_ENV=production
    depends_on:
      - tailscale
  tailscale:
    image: tailscale/tailscale:latest
    volumes:
      - ./tailscale:/var/lib/tailscale
    environment:
      - TS_AUTH_KEY=${TS_AUTH_KEY}
      - TS_HOSTNAME=app-server
      - TS_EXTRA_ARGS=--advertise-exit-node
    cap_add:
      - NET_ADMIN
      - SYS_MODULE
    restart: unless-stopped

```

### 3. Tailscale Configuration File

```

JSON
// tailscale.json
{
  "hostname": "repo-dev-environment",
  "acls": [
    {
      "action": "accept",
      "users": ["group:developers"],
      "ports": ["*:*"]
    }
  ],
  "tags": ["env:development"],
  "advertised_routes": ["10.0.0.0/24"],
  "exit_node": false
}

```

### 4. GitHub Repository Settings

Add these secrets to your GitHub repository:

- `TS_AUTH_KEY`: Your Tailscale authentication key
- `TS_OAUTH_CLIENT_ID`: OAuth Client ID for GitHub Actions
- `TS_OAUTH_SECRET`: OAuth Secret for GitHub Actions
- `DOTENV_KEY`: Your dotenv-vault key



## 5. Combined Implementation for CI/CD

```
YAML

# .github/workflows/deploy-secure.yml
name: Secure Deploy
on:
  push:
    branches: [main]
jobs:
  deploy:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3

      - name: Setup Node.js
        uses: actions/setup-node@v3
        with:
          node-version: '18'

      - name: Setup Tailscale
        uses: tailscale/github-action@v2
        with:
          oauth-client-id: ${ secrets.TS_OAUTH_CLIENT_ID }
          oauth-secret: ${ secrets.TS_OAUTH_SECRET }
          tags: tag:ci

      - name: Load environment variables
        run: npx dotenv-vault@latest pull production
        env:
          DOTENV_KEY: ${ secrets.DOTENV_KEY }

      - name: Deploy to secure infrastructure
        run: |
          # Now you can access internal resources via Tailscale
          # with secure environment variables loaded
          npm run deploy:secure
```

## 6. Security Best Practices

### 1. Rotate Keys Regularly

- Set up automated rotation for Tailscale and dotenv-vault keys

### 2. Access Control

- Use Tailscale ACLs to limit which users can access which services
- Configure dotenv-vault with granular access controls for team members

### 3. Auditing

- Enable Tailscale audit logs
- Track dotenv-vault access with [dotenv-vault access](#)

#### 4. Documentation

- Create / Update `SECURITY.md` file explaining your security approach
- Document the process for requesting access to environments

This implementation combines the security benefits of encrypted environment variables with dotenv-vault and secure network access with Tailscale, creating a robust security posture for your repositories and infrastructure.

---

## Resource Gathering

### ▼ ○ 5 Agent Frameworks & 3 Protocols Frameworks

#### ▼ ○ AG2



AG2 - ag2

<https://docs.ag2.ai/latest/>



GitHub - ag2ai/fastagency: The fastest way to bring multi-agent workflows to production.

<https://github.com/ag2ai/fastagency?tab=readme-ov-file>

#### ▼ ○ Google ADK



Agent Development Kit

<https://google.github.io/adk-docs/>

#### ▼ ○ AutoGPT



AutoGPT Documentation

<https://dev-docs.agpt.co/>



GitHub - Significant-Gravitas/AutoGPT: AutoGPT is the vision of accessible AI for everyone, to use a...

<https://github.com/Significant-Gravitas/AutoGPT>

#### ▼ ○ Agno



What is Agno - Agno

<https://docs.agno.com/introduction>

#### ▼ ○ Lang [\_\_\_\_\_]



[LangChain](#)



[LangGraph](#)




[Langflow | Low-code AI builder for agentic and RAG applications](#)








[Langfuse](#)



[LangSmith](#)

- ▶ ☐ [Introduction - LangWatch](#)
- ▼ ☐ Open Agent Platform
  -  Quick Start - Open Agent Platform  
<https://docs.oap.langchain.com/quickstart>
- ▶ ☐ [Just a moment...](#)

## Protocols

- ▼ ☐ A2A (google)
    -  Agent2Agent Protocol  
<https://google.github.io/A2A/#/>
  - ▼ ☐ Agent Protocol (Langchain)
    -  GitHub - langchain-ai/agent-protocol  
<https://github.com/langchain-ai/agent-protocol>
  - ▼ ☐ Agent Protocol (AutoGPT)
    -  GitHub - Significant-Gravitas/agent-protocol: Common interface for interacting with AI agents. The p...  
<https://github.com/Significant-Gravitas/agent-protocol>
    -  Agent Protocol  
<https://agentprotocol.ai/>
  - ▼ ☐ Agent IDEs
    - ▼ ☐ AG2 Studio
      -  GitHub - ag2ai/ag2studio: AG2 Studio  
<https://github.com/ag2ai/ag2studio>
    - ▼ ☐ LangGraph Studio
      - ☐ Overview  
[https://langchain-ai.github.io/langgraph/concepts/langgraph\\_studio/](https://langchain-ai.github.io/langgraph/concepts/langgraph_studio/)
    - ▼ ☐ Letta ADE
      - ☐ Letta  
<https://app.letta.com/development-servers/local/agents/agent-a44fb81e-fd9c-48ed-b119-31ab219b5567>
    - ▼ ☐ Letta Docs
      - ☐ Agent Development Environment (ADE) | Letta  
<https://docs.letta.com/guides/ade/overview>
-

## ▼ ++ Agentic

### ▼ ○ Inngest



Inngest - AI and backend workflows, orchestrated at any scale  
<https://www.inngest.com/>



GitHub - inngest/inngest: The leading workflow orchestration platform. Run stateful step functions a...  
<https://github.com/inngest/inngest>

### ▼ ○ Inngest Agent Kit



GitHub - inngest/agent-kit: AgentKit: Build multi-agent networks in TypeScript with deterministic ro...  
<https://github.com/inngest/agent-kit>

## Agent Kit Reddit Search BrowserBase



agent-kit/examples/reddit-search-browserbase-tools at main · inngest/agent-kit  
<https://github.com/inngest/agent-kit/tree/main/examples/reddit-search-browserbase-tools>

---

## No Code / Low Code

### Flowize AI



Flowize - Build AI Agents, Visually  
<https://flowiseai.com/>



GitHub - FlowiseAI/Flowise: Build AI Agents, Visually  
<https://github.com/FlowiseAI/Flowise>

### Langflow



Langflow | Low-code AI builder for agentic and RAG applications  
<https://www.langflow.org/>

---

### ▼ ○ Scrape & Crawl



GitHub - unclecode/crawl4ai: 🕸️🤖 Crawl4AI: Open-source LLM Friendly Web Crawler & Scraper. Don't be...  
<https://github.com/unclecode/crawl4ai>

## LightRAG



GitHub - HKUDS/LightRAG: "LightRAG: Simple and Fast Retrieval-Augmented Generation"  
<https://github.com/HKUDS/LightRAG>

---

## ▼ Browser Control

## Browser Use



Browser Use Agents & Tools



Resource

### ▼ Autonomous

#### Browser User



Introduction - Browser Use

<https://docs.browser-use.com/introduction>



GitHub - browser-use/web-ui: Run AI Agent in your browser.

<https://github.com/browser-use/web-ui>



mcp-browser-use/src/mcp\_server\_browser\_use/server.py at main · Saik0s/mcp-browser-use

[https://github.com/Saik0s/mcp-browser-use/blob/main/src/mcp\\_server\\_browser\\_use/server.py](https://github.com/Saik0s/mcp-browser-use/blob/main/src/mcp_server_browser_use/server.py)

#### Proxy Lite



GitHub - convergence-ai/proxy-lite: A mini, open-weights, version of our Proxy assistant.

<https://github.com/convergence-ai/proxy-lite/tree/main>

#### Open Manus



GitHub - mannaandpoem/OpenManus: No fortress, purely open ground. OpenManus is Coming.

<https://github.com/mannaandpoem/OpenManus>

#### Open Operator



GitHub - browserbase/open-operator

<https://github.com/browserbase/open-operator>

#### Steel



Steel | Open-source Headless Browser API

<https://steel.dev/>

#### BrowserBase



Browserbase - Headless Web Browser API

<https://www.browserbase.com/overview>

### ▼ Production



Cloudflare

☐ Welcome to Cloudflare  
<https://developers.cloudflare.com/>

☐ DNS

☐ Move Domain Name to CF

☐ Security / WAF

☐ Deploy Managed Rules

☐ Deploy Robots.txt

---

▼ **Stack As TODOs**

▼ **(Inside) DevAgents & DevTools**

☒ Agent Zero

☒ Task

☒ BrainCraft

☒ Task

☒ GPT4Free

☒ Task

☒ Aider

☒ Task

☒ Genaroo

☒ Task

▼ **System Architecture**

☒ Capture Remaining Docs

☒ Task

☒ Synthesize Into PRD & Description for GenA-Roo

☒ Task

▼ **Database**

☒ SurrealDB & Surrealist

☒ Task

**SpacetimeDB**

☐ SpacetimeDB  
<https://spacetimedb.com/docs>


**Neon & Turso**

☐ Neon and Turso for Penumbra

☐ Page


**Neon**

☐ Neon Serverless Postgres — Ship faster  
<https://neon.tech/>

 GitHub - neondatabase/neon: Neon: Serverless Postgres. We separated storage and compute to offer aut...  
<https://github.com/neondatabase/neon>

## Turso

☐ Turso - SQLite Databases for all Apps  
<https://turso.tech/>

 Welcome to Turso Cloud - Turso  
<https://docs.turso.tech/introduction>

## Memory

☒ Mem0

☒ Task

☒ Letta (MemGPT)

☒ Task


☒ MemGraph


☒ Task


## ▼ GenUI

### Crayon

☐ Introduction | CrayonAI  
<https://crayonai.org/docs/>

 GitHub - thesysdev/crayon: Generative UI SDK  
<https://github.com/thesysdev/crayon>

 examples/crayon at main · thesysdev/examples  
<https://github.com/thesysdev/examples/tree/main/crayon>

 @storybook/core - Storybook  
<https://crayonai.org/ui/?path=/docs/components-accordion--docs>

## ▼ Infra & IaC

☒ Temporal

☒ Task

## ▼ DevSecOps & CI/CD

### ▼ Docker

☐ Docker Engine  
<https://docs.docker.com/engine/>

### MCP Catalog & Toolkit

☐ MCP Catalog and Toolkit

<https://docs.docker.com/ai/mcp-catalog-and-toolkit/>

## GH Copilot

- ☐ Docker for GitHub Copilot  
<https://docs.docker.com/copilot/>

## Docker Compose

- ☐ Docker Compose  
<https://docs.docker.com/compose/>

## Docker Bake

- ☐ Bake  
<https://docs.docker.com/build/bake/>

## Testcontainers Cloud

- ☐ Docs for Testcontainers Cloud  
<https://testcontainers.com/cloud/docs/>


### ▼ k8s

- ☐ Production-Grade Container Orchestration  
<https://kubernetes.io/>

### ▼ Dagger


- ☐ Dagger.io  
<https://dagger.io/>

### ▼ Dagger Agents

- ☐  Dagger Agents

 Page

### ▼ Uffizzi

- ☒  Uffizzi – Uffizzi  
<https://docs.uffizzi.com/>
- ☐ Uffizzi app  
[https://app.uffizzi.com/account\\_settings#integrations](https://app.uffizzi.com/account_settings#integrations)

## Temporal

- ☐ Temporal Platform Documentation | Temporal Platform Documentation  
<https://docs.temporal.io/>

## GitHub Actions

- ☐ GitHub Actions  
<https://github.com/features/actions>
- ☐ GitHub Actions documentation - GitHub Docs



<https://docs.github.com/en/actions>

## GitHub Dev Program

- ☐ GitHub Developer Program - GitHub Docs  
<https://docs.github.com/en/get-started/exploring-integrations/github-developer-program>


## ▼ LLMetry & Ops

### ▼ Promptfoo

- ☐ Promptfoo - Portkey Docs  
<https://portkey.ai/docs/integrations/libraries/promptfoo>

### ▼ Comet / Opik

#### Opik MCP


-  GitHub - comet-ml/opik-mcp: Model Context Protocol (MCP) implementation for Opik enabling seamless I...  
<https://github.com/comet-ml/opik-mcp>

#### Opik

- Debug, trace, evaluate and monitor LLM apps and RAG pipelines

- ☐ Comet Opik  
<https://www.comet.com/opik/areid987/home>

#### Opik GitHub

-  GitHub - comet-ml/opik: Debug, evaluate, and monitor your LLM applications, RAG systems, and agentic...  
<https://github.com/comet-ml/opik>

#### Comet Docs

- ☐ Docs Home - Comet Docs  
<https://www.comet.com/docs/v2/>

#### Comet GitHub

-  Comet  
<https://github.com/comet-ml>

### ▼ AgentOps

- ☐ AgentOps  
<https://www.agentops.ai/>

### ▼ Agenta

- ☐ What is Agenta? - Docs - Agenta

<https://docs.agenta.ai/>



GitHub - Agenta-AI/agenta: The open-source LLMOps platform: prompt playground, prompt management, LL...  
<https://github.com/agenta-ai/agenta>

## ▼ Langwatch



Introduction - LangWatch  
<https://docs.langwatch.ai/introduction>

## ▼ Service Mesh

### Kong API & AI

#### Kuma



Kuma Service Mesh

Page

#### Kong Gateway



Installation Options - Kong Gateway | Kong Docs  
<https://docs.konghq.com/gateway/latest/install/#kong-community>

#### Portkey



What is Portkey? - Portkey Docs  
<https://portkey.ai/docs/introduction/what-is-portkey>

#### Not Diamond



What is Not Diamond?  
<https://docs.notdiamond.ai/docs/what-is-not-diamond>

## ▼ LiteLLM



LiteLLM  
<https://www.litellm.ai/>



Docker, Deployment | LiteLLM  
<https://docs.litellm.ai/docs/proxy/deploy>

---

## ▼ CMS

### Prismic




Developer Documentation - Prismic  
<https://prismic.io/docs>

---

## ▼ More Agentic Frameworks

## Dagger

- ☐ Dagger Documentation | Dagger  
<https://docs.dagger.io/>
-  GitHub - dagger/agents  
<https://github.com/dagger/agents>

## Langbase ICP Insight Team Blogger

- ☐ icp-insight-team-agent-0677 - areid987  
<https://chai.new/areid987/icp-insight-team-agent-0677>


## AI Devin

- ☐ Build a composable AI Devin - Guides  
<https://langbase.com/docs/guides/build-composable-ai-devin>
- ☐ DataStax Langflow | DataStax Docs  
<https://docs.datastax.com/en/langflow/index.html>
- ☐ Datastax Astra  
<https://astra.datastax.com/org/67b88500-04ee-48b8-a773-eabdc1459116>
- ☐ Untitled  
<https://oap.langchain.com/?agentId=eaad1967-3353-4897-a053-3d6ee59dffbe&deploymentId=ee18b042-9d0a-4994-8920-f8af1604a454&threadId=84e5b5de-1220-4603-932a-bf44d23a0d01>
- ☐ Code Alchemist - Write code or design database with AI  
<https://code-alchemist.langbase.dev/>
- ☐ Agentic Generative UI  
<https://docs.copilotkit.ai/coagents/generative-ui/agentic>


---

## ML Ops

### ▼ Prime

-  GitHub - PrimeIntellect-ai/prime: prime is a framework for efficient, globally distributed training ...  
<https://github.com/PrimeIntellect-ai/prime>

### Q Learn

-  GitHub - Div99/IQ-Learn: (NeurIPS '21 Spotlight) IQ-Learn: Inverse Q-Learning for Imitation  
<https://github.com/Div99/IQ-Learn>

### ▼ Graph Resources

-  Graph Resources

 Resource